

HOUSING PRICE PREDICTION

Platform – ANACONDA
Programming Language – PYTHON

HARSH YADAV

HARSH YADAV

INDEX

- OBJECTIVE
- IMPORTING LIBRARIES & LOADING DATASET
- EXPLORING DATA
- LINEAR REGRESSION (WITHOUT TREATING OUTLIERS)
- LINEAR REGRESSION (AFTER TREATING OUTLIERS)
- DECISION TREE REGRESSION
- RANDOM FOREST REGRESSION
- XGBOOST
- GRADIENT BOOSTING REGRESSION
- CROSS-VALIDATION SCORES

OBJECTIVE

- A Real estate company that runs a website for buying and selling of houses, they would like to give the customers a prediction of the cost of the houses in a few years.
- We have to use the housing data, we have to pre-process it, and have to create a model that will predict the price of the house and will measure the accuracy of the model.



IMPORTING LIBRARIES

Housing Price Prediction

(using LinearRegression Model)

- HARSH YADAV

In []:

IMPORTING LIBRARIES & DATASET

```
In [1]: import pandas as pd          # to import dataset and create data frames
       import numpy as np           # for solving complex mathematical problems
       import matplotlib.pyplot as plt # tool for data visualization & representation
       import seaborn as sns # based on matplotlib, but with more colourful themes
       import sklearn

       import statsmodels.api as smapi
       from sklearn.linear_model import LinearRegression
       # NOTE: we can use any of the above library for linear regression model,
       #        but i am using scikit_learn because its best for big data (no. of observations > 1000)

In [2]: data = pd.read_csv(r"C:\Users\anika\OneDrive\Desktop\data_science\Housing Price Prediction.csv")
```

HARSH YADAV

EXPLORING DATA



HARSH YADAV

EXPLORING DATA

```
In [3]: data.head()
Out[3]:
      id      date    price  bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront  view ... grade  sqft_above  sqft_basement  yr_built
0  9106000005  5/27/2015  1310000.0       4     2.25       3750     5000     2.0        0     0 ...     8     2440           0     1310     19
1  5101400871  5/24/2015  445500.0       2     1.75      1390     6670     1.0        0     0 ...     6     720           0     670     19
2  7923600250  5/15/2015  450000.0       5     2.00      1870     7344     1.5        0     0 ...     7     1870           0     0     19
3  8730000270  5/14/2015  359000.0       2     2.75      1370     1140     2.0        0     0 ...     8     1080           0     290     20
4  9178601660  5/14/2015  1700000.0       5     3.00      3320     5354     2.0        0     0 ...     9     3320           0     0     20
5 rows × 21 columns
```

```
In [4]: data.describe()
Out[4]:
      id      price  bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront  view  condition
count  2.159700e+04  2.159700e+04  21597.000000  21597.000000  2.159700e+04  21597.000000  21597.000000  21597.000000  21597.000000
mean   4.580474e+09  5.402966e+05   3.373200   2.115826  2080.321850  1.509941e+04   1.494096   0.007547   0.234292   3.409825
std    2.876736e+09  3.673681e+05   0.926299   0.768984  918.106125  4.141264e+04   0.539683   0.086549   0.766390   0.650546
min    1.000102e+06  7.800000e+04   1.000000   0.500000  370.000000  5.200000e+02   1.000000   0.000000   0.000000   1.000000
25%   2.123049e+09  3.220000e+05   3.000000   1.750000  1430.000000  5.040000e+03   1.000000   0.000000   0.000000   3.000000
50%   3.904930e+09  4.500000e+05   3.000000   2.250000  1910.000000  7.618000e+03   1.500000   0.000000   0.000000   3.000000
75%   7.308900e+09  6.450000e+05   4.000000   2.500000  2550.000000  1.068500e+04   2.000000   0.000000   0.000000   4.000000
max   9.900000e+09  7.700000e+06   33.000000   8.000000  13540.000000  1.651359e+06   3.500000   1.000000   4.000000   5.000000
```

```
In [5]: data.info()
Out[5]:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
 #   Column          Non-Null Count  Dtype  
--- 
 0   id              21597 non-null   int64  
 1   date            21597 non-null   object 
 2   price           21597 non-null   float64 
 3   bedrooms        21597 non-null   int64  
 4   bathrooms       21597 non-null   float64 
 5   sqft_living     21597 non-null   int64  
 6   sqft_lot        21597 non-null   int64  
 7   floors          21597 non-null   float64 
 8   waterfront      21597 non-null   int64  
 9   view            21597 non-null   int64  
 10  condition       21597 non-null   int64  
 11  grade           21597 non-null   int64  
 12  sqft_above      21597 non-null   int64  
 13  sqft_basement   21597 non-null   int64  
 14  yr_builtin      21597 non-null   int64  
 15  yr_renovated    21597 non-null   int64  
 16  zipcode         21597 non-null   int64  
 17  lat             21597 non-null   float64 
 18  long            21597 non-null   float64 
 19  sqft_living15   21597 non-null   int64  
 20  sqft_lot15      21597 non-null   int64  
dtypes: float64(5), int64(15), object(1)
memory usage: 3.5+ MB
```

```
In [6]: data.corr() # TO GET CORELLATION BETWEEN EVERY FEATURE OF DATASET
```

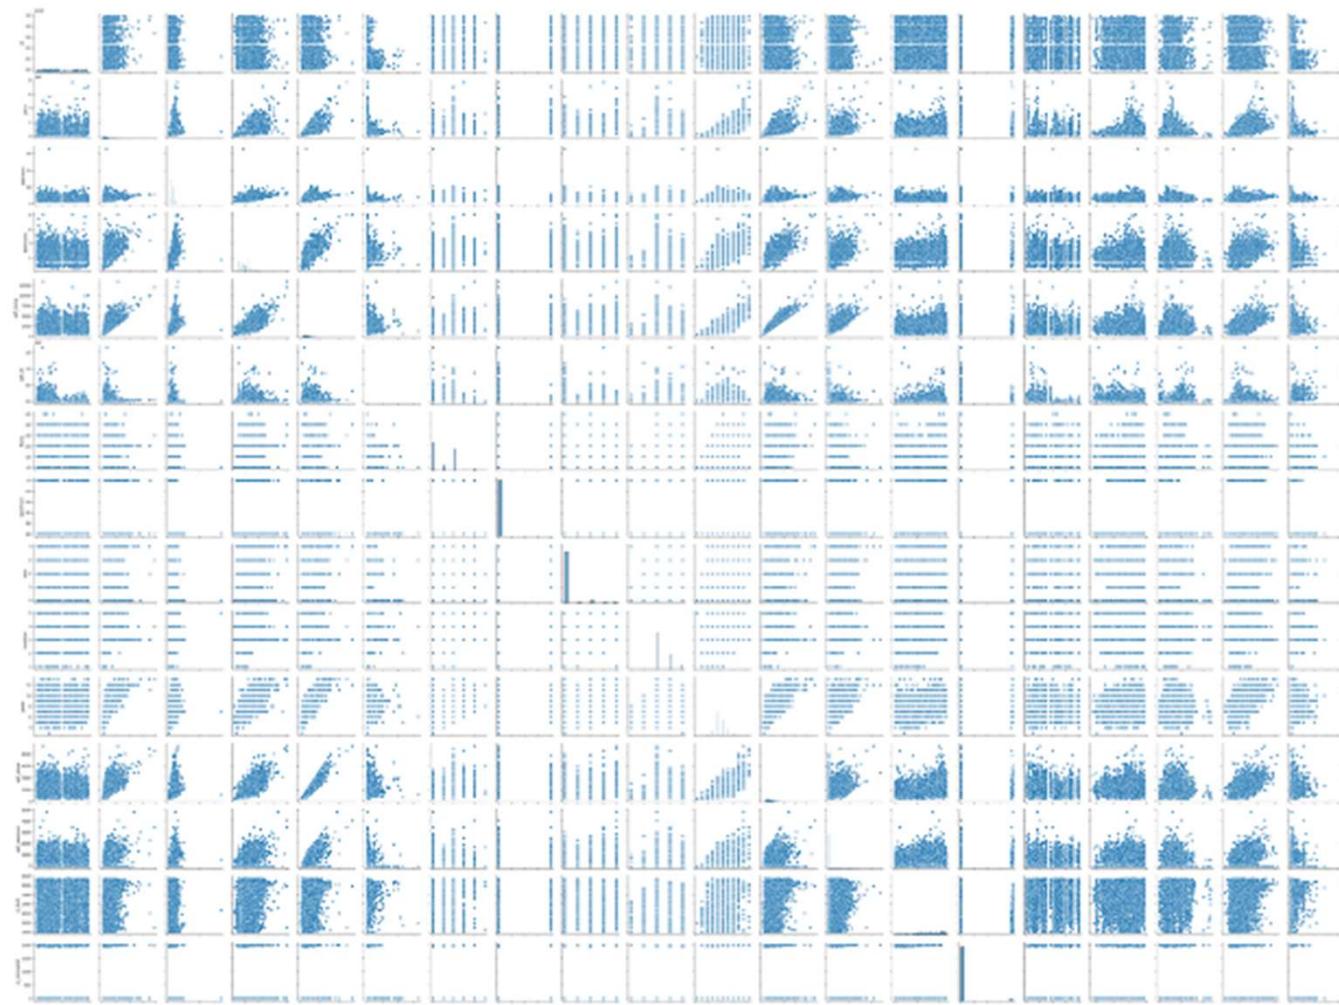
Out[6]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	sqft_lot15
id	1.000000	-0.016772	0.001150	0.005162	-0.012241	-0.131911	0.018608	-0.002727	0.011536	-0.023803	0.008188	-0.010799								
price	-0.016772	1.000000	0.308787	0.525906	0.701917	0.089876	0.256804	0.266398	0.397370	0.036056	0.667951	0.605368								
bedrooms	0.001150	0.308787	1.000000	0.514508	0.578212	0.032471	0.177944	-0.006834	0.080008	0.026496	0.356563	0.479386								
bathrooms	0.005162	0.525906	0.514508	1.000000	0.755758	0.088373	0.502582	0.063744	0.188386	-0.126479	0.665838	0.686668								
sqft_living	-0.012241	0.701917	0.578212	0.755758	1.000000	0.173453	0.353953	0.103854	0.284709	-0.059445	0.762779	0.876448								
sqft_lot	-0.131911	0.089876	0.032471	0.088373	0.173453	1.000000	-0.004814	0.021632	0.074900	-0.008830	0.114731	0.184139								
floors	0.018608	0.256804	0.177944	0.502582	0.353953	-0.004814	1.000000	0.023755	0.028814	-0.264075	0.458794	0.523989								
waterfront	-0.002727	0.266398	-0.006834	0.063744	0.103854	0.021632	0.023755	1.000000	0.401971	0.016611	0.082888	0.072109								
view	0.011536	0.397370	0.080008	0.188386	0.284709	0.074900	0.028814	0.401971	1.000000	0.045999	0.251728	0.167609								
condition	-0.023803	0.036056	0.026496	-0.126479	-0.059445	-0.008830	-0.264075	0.016611	0.045999	1.000000	-0.146896	-0.158904								
grade	0.008188	0.667951	0.356563	0.665838	0.762779	0.114731	0.458794	0.082888	0.251728	-0.146896	1.000000	0.756073								
sqft_above	-0.010799	0.605368	0.479386	0.686668	0.876448	0.184139	0.523989	0.072109	0.167609	-0.158904	0.756073	1.000000								
sqft_basement	-0.005193	0.323799	0.302808	0.283440	0.435130	0.015418	-0.245715	0.080559	0.277078	0.173849	0.168220	-0.052156								
yr_built	0.021617	0.053953	0.155670	0.507173	0.318152	0.052946	0.489193	-0.026153	-0.053636	-0.361592	0.447865	0.424037								
yr_renovated	-0.016925	0.126424	0.018389	0.050544	0.055308	0.007686	0.006427	0.092873	0.103951	-0.060788	0.014261	0.023251								
zipcode	-0.008211	-0.053402	-0.154092	-0.204786	-0.199802	-0.129586	-0.059541	0.030272	0.084622	0.002888	-0.185771	-0.261570								
lat	-0.001798	0.306692	-0.009951	0.024280	0.052155	-0.085514	0.049239	-0.014306	0.005871	-0.015102	0.113575	-0.001199								
long	0.020672	0.022036	0.132054	0.224903	0.241214	0.230227	0.125943	-0.041904	-0.078107	-0.105877	0.200341	0.344842								
sqft_living15	-0.002701	0.585241	0.393406	0.569884	0.756402	0.144763	0.280102	0.086507	0.280681	-0.093072	0.713867	0.731767								
sqft_lot15	-0.138557	0.082845	0.030690	0.088303	0.184342	0.718204	-0.010722	0.030781	0.072904	-0.003126	0.120981	0.195077								

HARSH YADAV

```
In [92]: sns.pairplot(data)
```

```
Out[92]: <seaborn.axisgrid.PairGrid at 0x1ef40353850>
```



HARSH YADAV

checking for missing_values

In [8]: `data.isnull().sum()`

```
Out[8]: id      0  
date     0  
price    0  
bedrooms 0  
bathrooms 0  
sqft_living 0  
sqft_lot   0  
floors    0  
waterfront 0  
view      0  
condition  0  
grade     0  
sqft_above 0  
sqft_basement 0  
yr_built   0  
yr_renovated 0  
zipcode    0  
lat       0  
long      0  
... 0
```

In [9]: `#no null values are there`

```
In [10]: #No need to create dummy variables  
#just we have to deal with "date" as it is of data_type = "object"  
#there are a lot ways to perform with "High Cardinality Variables" like "Date"  
#we can:  
#1) - Embed with Freq  
#2) - Target Encoding  
#3) - or Simply Drop  
# I am using 3rd technique here,
```

```
In [11]: data.dtypes
```

```
Out[11]: id          int64  
date         object  
price        float64  
bedrooms     int64  
bathrooms    float64  
sqft_living   int64  
sqft_lot      int64  
floors        float64  
waterfront    int64  
view          int64  
condition     int64  
grade          int64  
sqft_above    int64  
sqft_basement int64  
yr_built      int64  
yr_renovated  int64  
zipcode       int64  
lat           float64  
long          float64  
sqft_living15 int64  
sqft_lot15    int64  
dtype: object
```

HARSH YADAV

Dropping the date variable, because of its "High Cardinality"

```
In [12]: data = data.drop(['date'],axis = 1)
```

```
In [13]: data.head()
```

Out[13]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_basement	yr_built
0	9106000005	1310000.0	4	2.25	3750	5000	2.0	0	0	5	8	2440		1310 1924
1	5101400871	445500.0	2	1.75	1390	6670	1.0	0	0	3	6	720		670 1941
2	7923600250	450000.0	5	2.00	1870	7344	1.5	0	0	3	7	1870		0 1960
3	8730000270	359000.0	2	2.75	1370	1140	2.0	0	0	3	8	1080		290 2009
4	9178601660	1700000.0	5	3.00	3320	5354	2.0	0	0	3	9	3320		0 2004

Removing targetted value "Price" from whole dataset "data"

```
In [14]: y = data.loc[:, "price"]
X = data.drop("price", axis = 1)
```

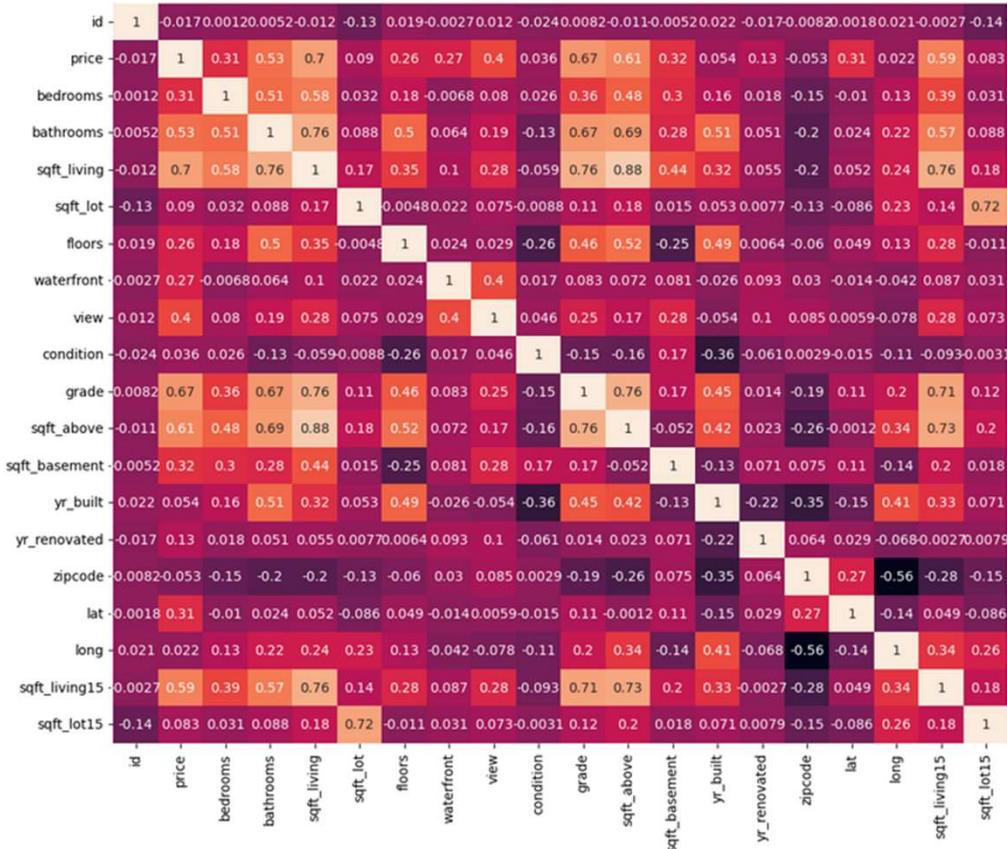
```
In [15]: X.head()
```

Out[15]:

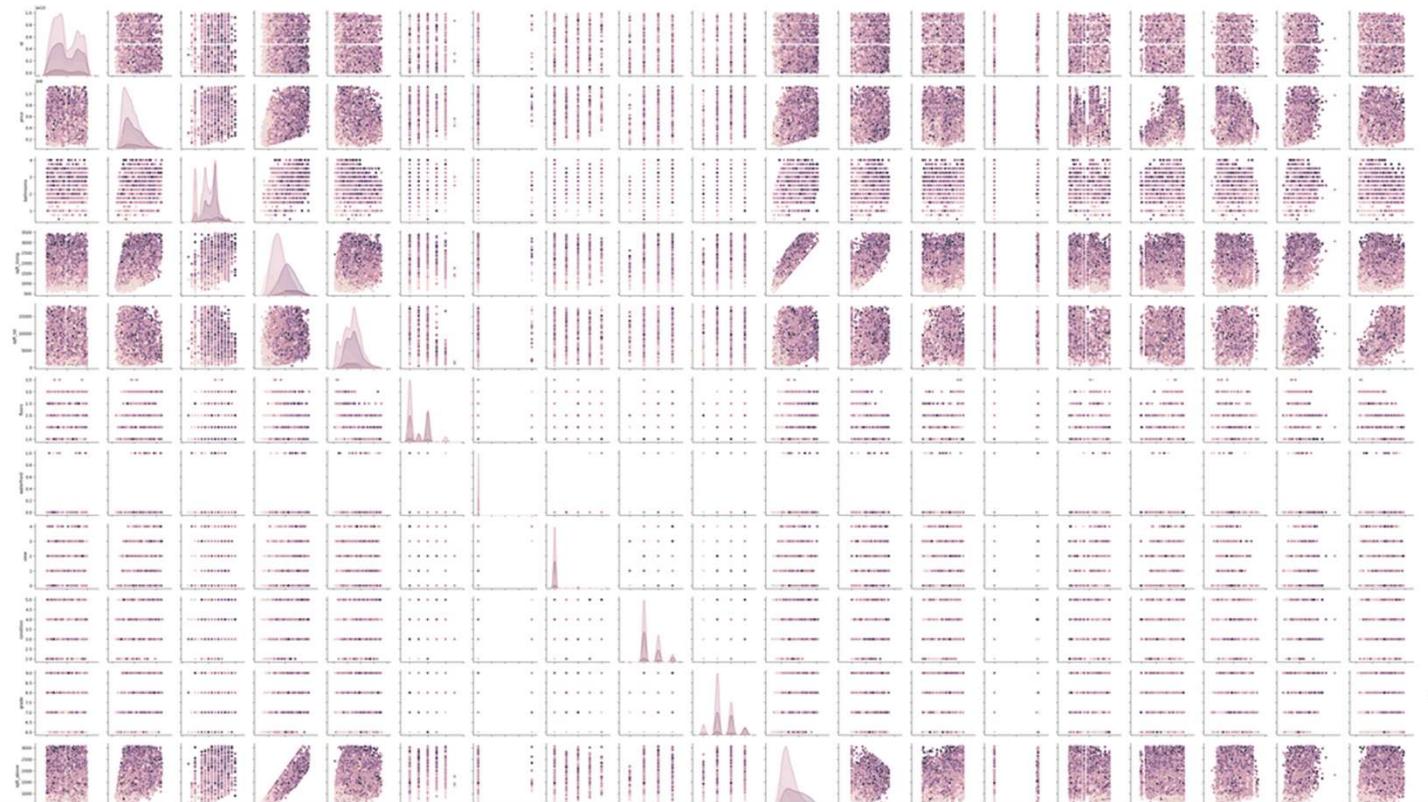
	id	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_basement	yr_built	yr_renovation
0	9106000005	4	2.25	3750	5000	2.0	0	0	5	8	2440		1310 1924	
1	5101400871	2	1.75	1390	6670	1.0	0	0	3	6	720		670 1941	
2	7923600250	5	2.00	1870	7344	1.5	0	0	3	7	1870		0 1960	
3	8730000270	2	2.75	1370	1140	2.0	0	0	3	8	1080		290 2009	
4	9178601660	5	3.00	3320	5354	2.0	0	0	3	9	3320		0 2004	

```
In [17]: plt.figure(figsize = (15,10))
sns.heatmap(data.corr(), annot = True)
# HEATMAP SHOWS US THE COLLINEARITY BY VISUALISATION & COLOUR GRADING
# WHICH MAKES IT REALLY EASY TO UNDERSTAND THE RELATION BETWEEN 2 VARIABLES
# HERE, LIGHT COLOURS ARE DEPICTING GREATER CORRELATION VALUE
# & DARK COLOURS ARE DEPICTING LESSER COORELATION VALUE
```

Out[17]: <AxesSubplot:>



```
In [89]: #by using 'hue' in pairplot, we can see how no. of bedrooms are corellated with other features of dataset  
#we can see the prices are low for light-colored dots, that means  
#dark ones are bhks like 33,10 &  
#light ones are bhk's like 2,3..  
sns.pairplot(df, hue = 'bedrooms')  
plt.savefig("hue")
```



```
In [90]: # PANDAS PROFILING REPORT WILL GIVE US THE DETAILED INFO ABOUT OUR DATA  
# THIS WILL HELP US UNDERSTAND OUR DATA, FEATURES, RECORDS, CORRELATIONS, MIN, MAX, COUNT AND A LOT MORE...  
from pandas_profiling import ProfileReport  
profile = ProfileReport(data, title="Pandas Profiling Report")
```

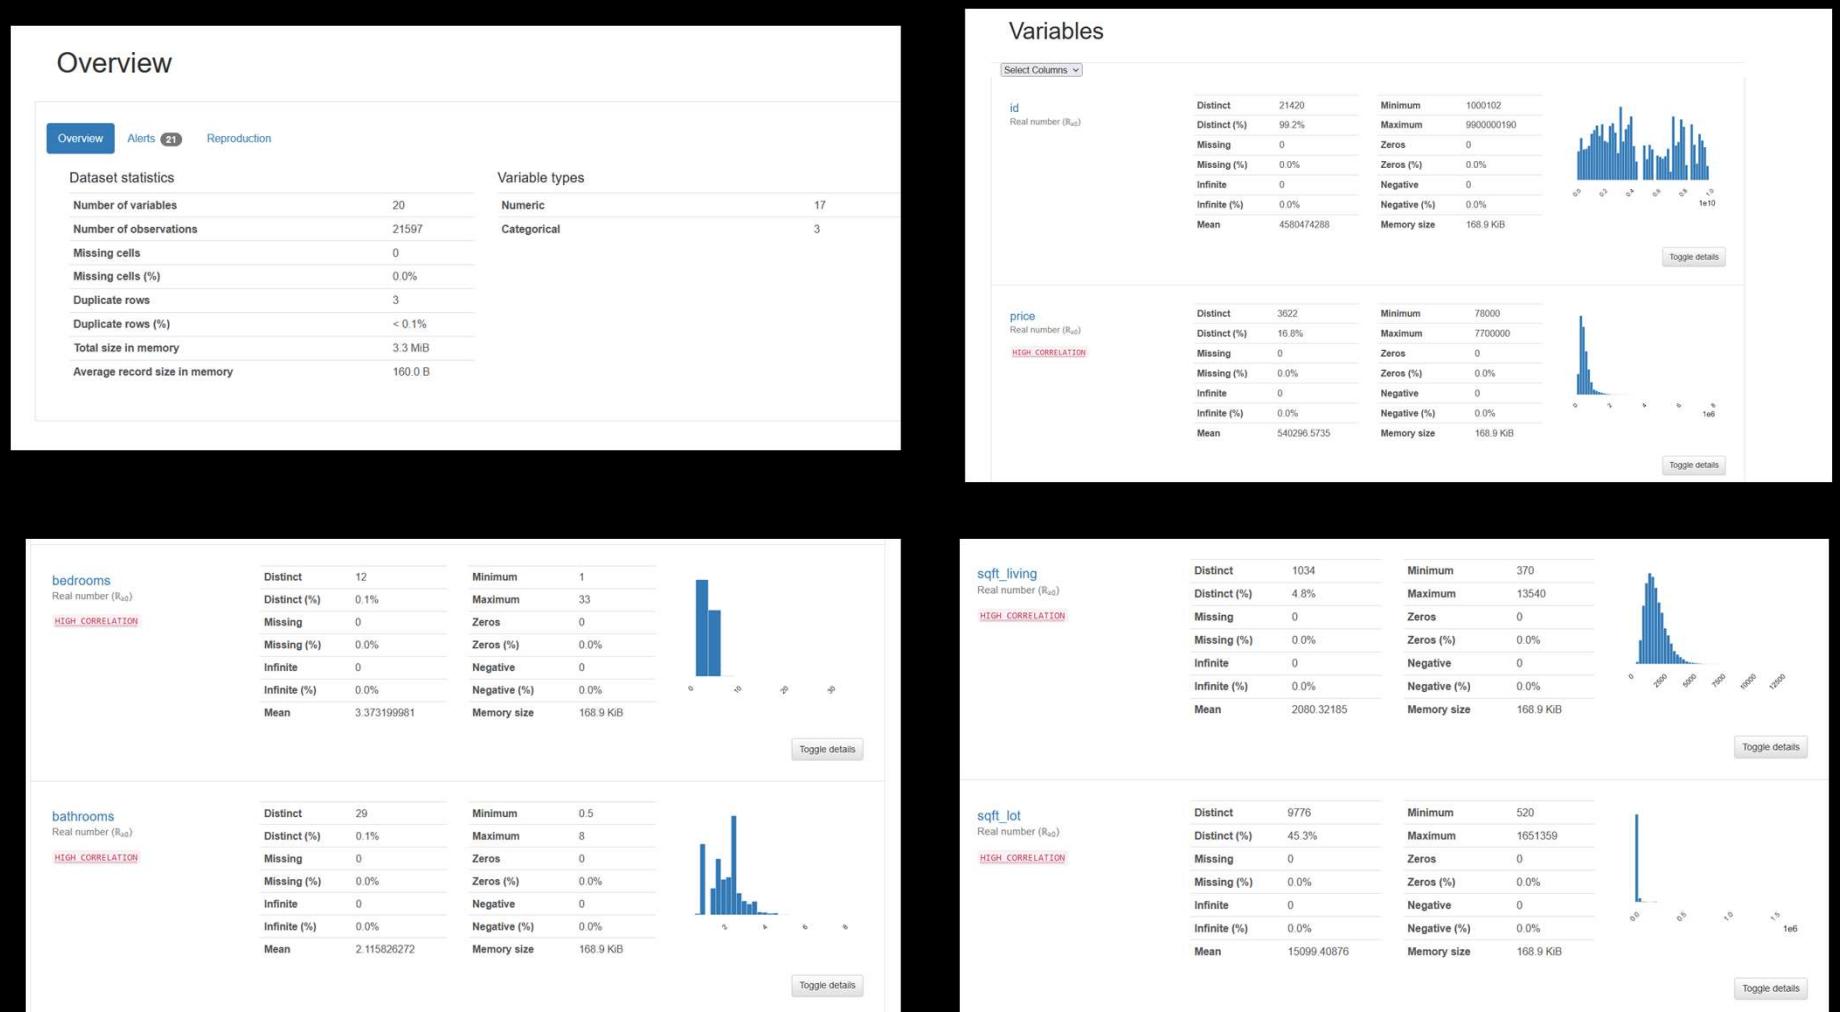
```
In [91]: profile.to_file("your_report.html")
```

Summarize dataset: 100%  323/323 [01:24<00:00, 3.44it/s, Completed]

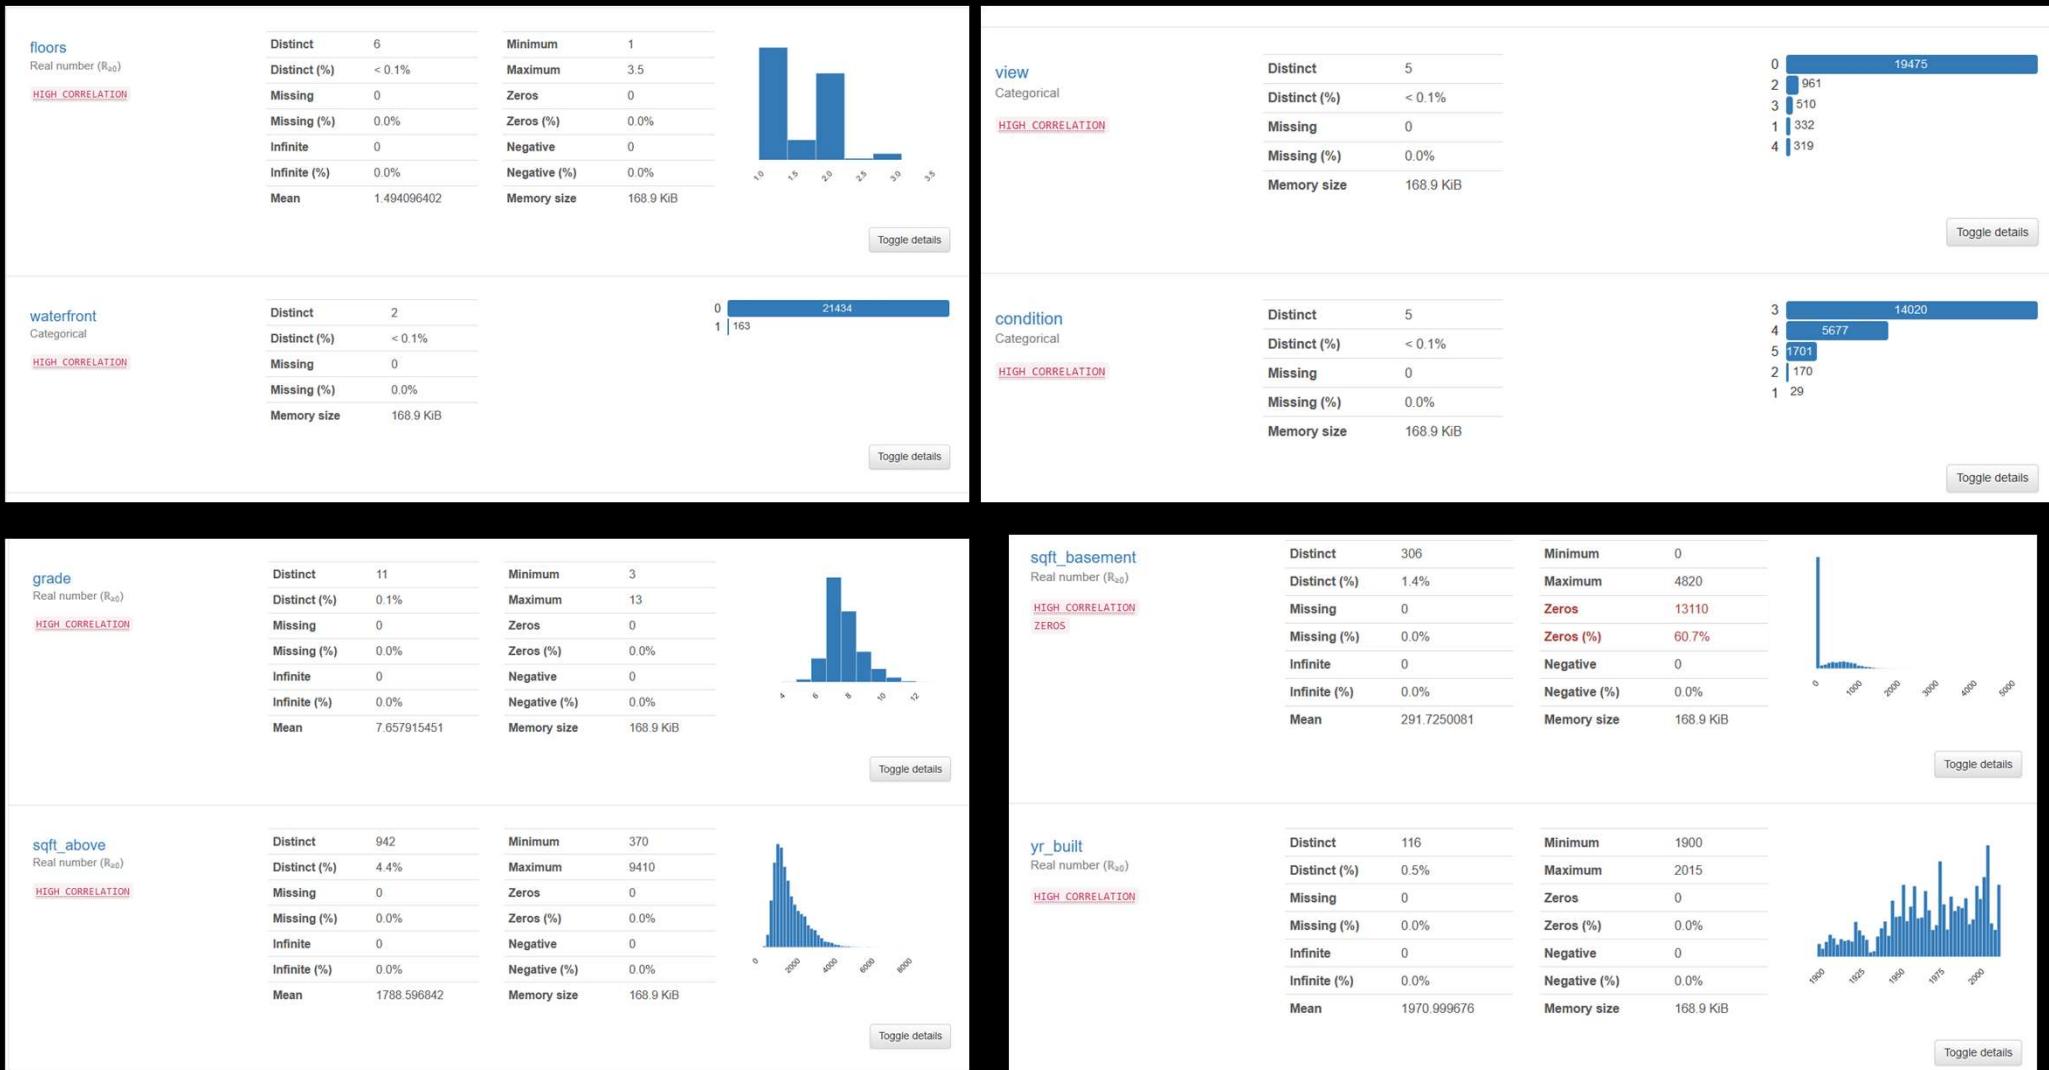
Generate report structure: 100%  1/1 [00:08<00:00, 8.23s/it]

Render HTML: 100%  1/1 [00:11<00:00, 11.29s/it]

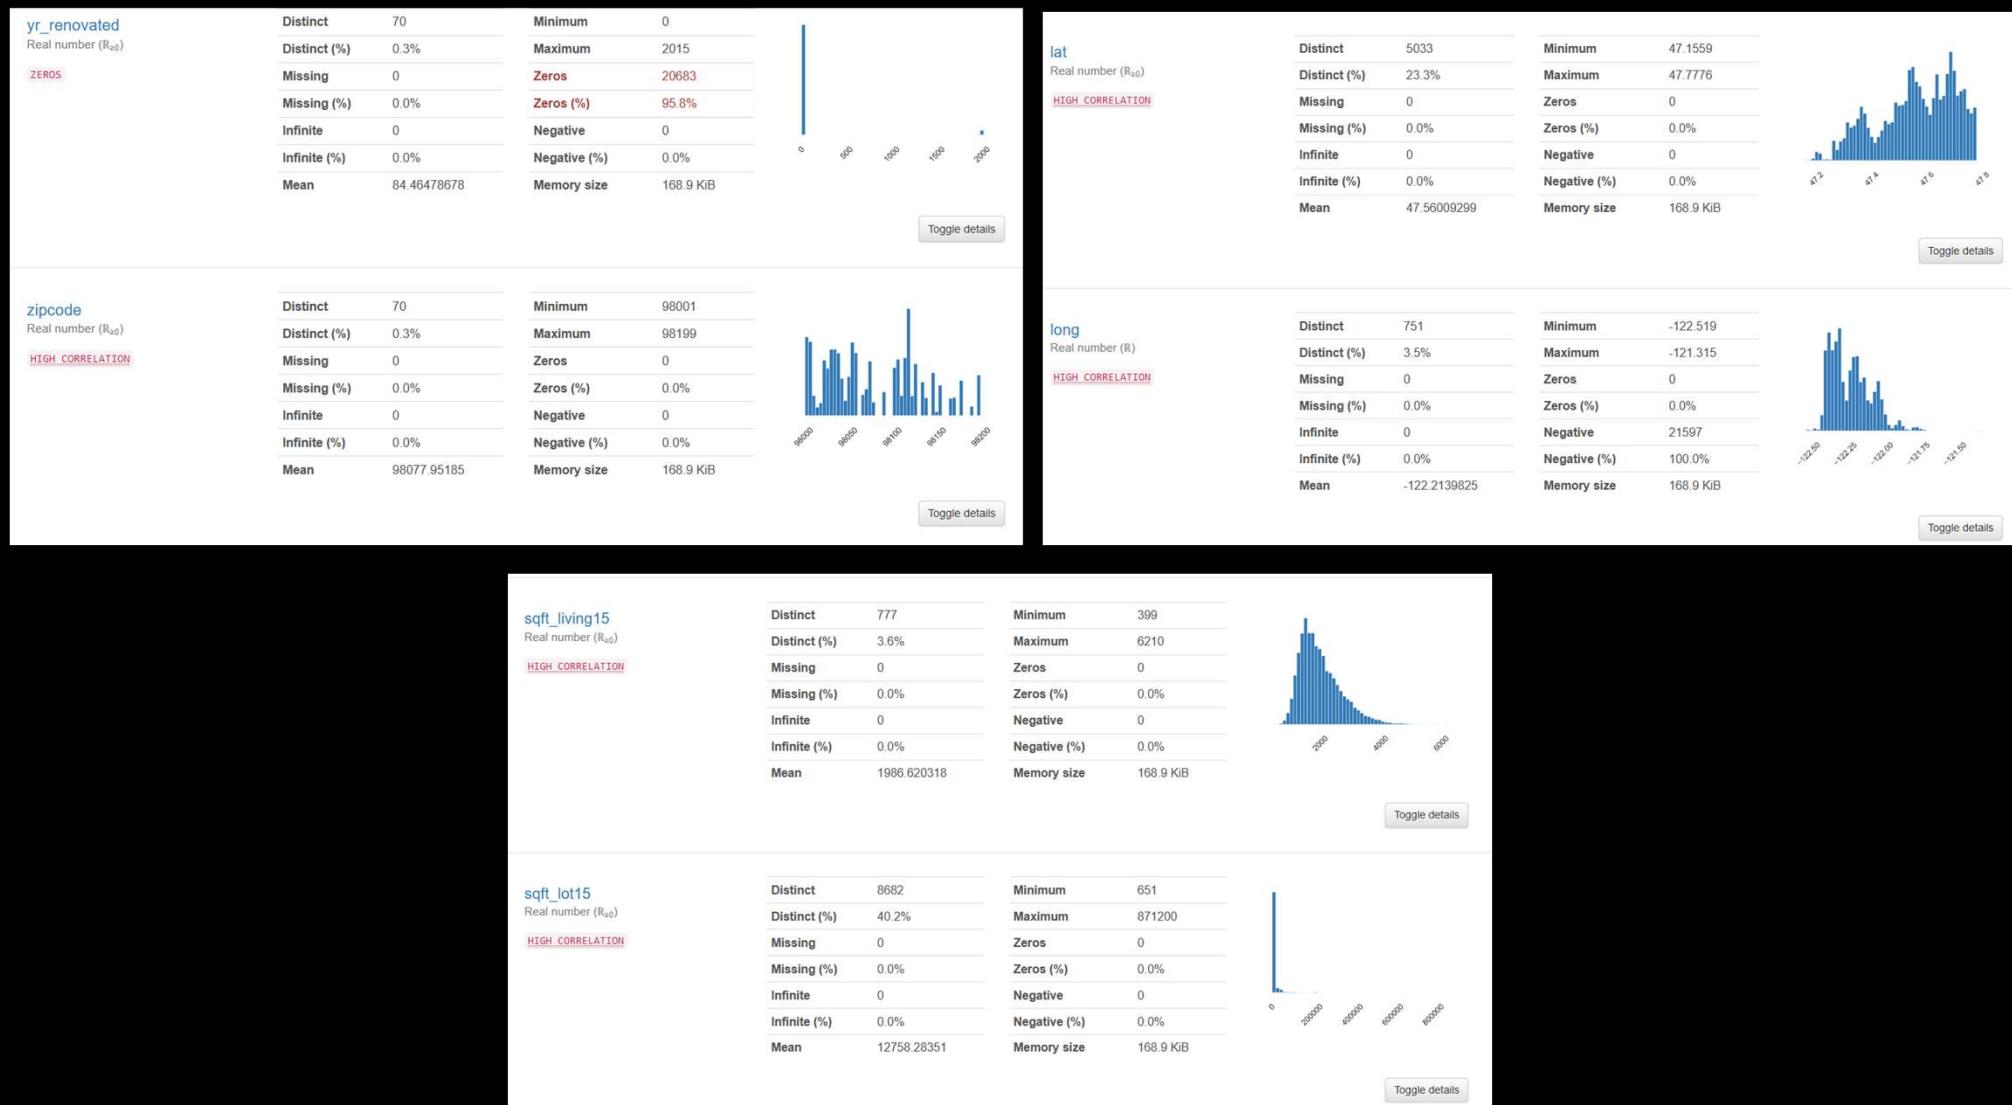
Export report to file: 100%  1/1 [00:00<00:00, 15.64it/s]



HARSH YADAV

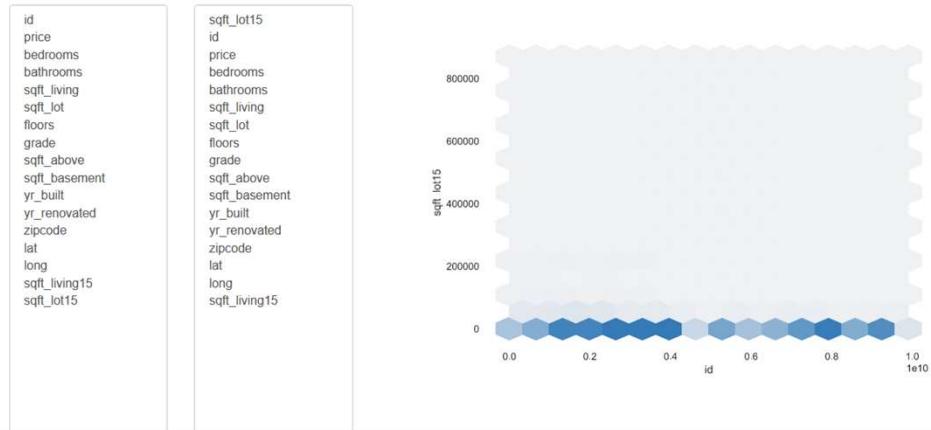


HARSH YADAV

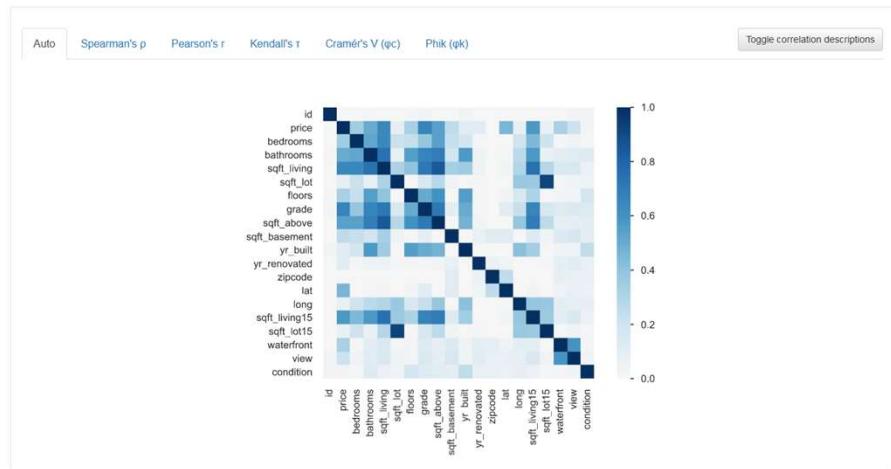


HARSH YADAV

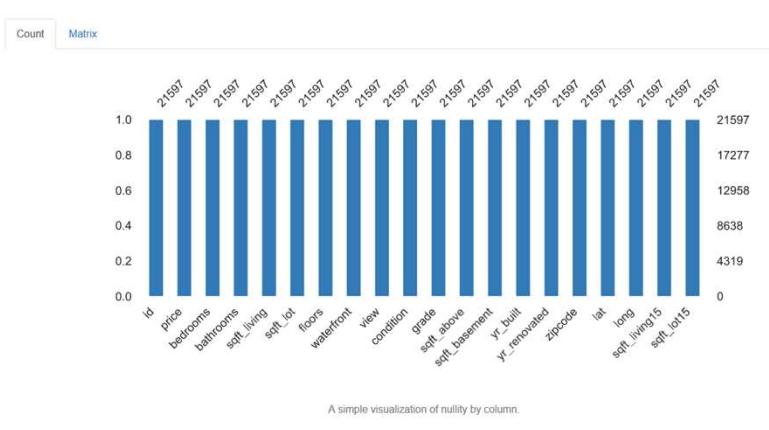
Interactions



Correlations



Missing values



Sample

First rows

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_basement	yr_built
0	9106000005	1310000.0	4	2.25	3750	5000	2.0	0	0	5	8	2440	1310	1924
1	5101400871	445500.0	2	1.75	1390	6670	1.0	0	0	3	6	720	670	1941
2	7923600250	450000.0	5	2.00	1870	7344	1.5	0	0	3	7	1870	0	1960
3	8730000270	359000.0	2	2.75	1370	1140	2.0	0	0	3	8	1080	290	2009
4	9178601660	1700000.0	5	3.00	3320	5354	2.0	0	0	3	9	3320	0	2004
5	1786200010	456500.0	4	2.50	2580	11780	2.0	0	0	3	9	2580	0	2003
6	1422700040	183000.0	3	1.00	1170	7320	1.0	0	0	3	7	1170	0	1962
7	6752600320	360000.0	4	2.50	2020	7289	2.0	0	0	3	7	2020	0	1994
8	4166600610	335000.0	3	2.00	1410	44866	1.0	0	0	4	7	1410	0	1985
9	7129304540	440000.0	5	2.00	1430	5600	1.5	0	0	3	6	1430	0	1947

HARSH YADAV

LINEAR REGRESSION (WITHOUT TREATING OUTLIERS)

Assumptions of Linear Regression



- 1 Linear Relationship
- 2 Normal Distribution of Residuals
- 3 Multicollinearity
- 4 Auto Correlation
- 5 Homoscedasticity

Last rows

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_basement	yr_built
21587	2201501015	430000.0	4	1.50	1920	10000	1.0	0	0	4	7	1070	850	1954
21588	2423600100	491500.0	4	1.75	2190	125452	1.0	0	2	3	9	2190	0	1968
21589	7853220390	785000.0	5	3.25	3660	11995	2.0	0	2	3	10	3660	0	2006
21590	9267200226	436110.0	3	2.50	1770	1235	3.0	0	0	3	8	1600	170	2007
21591	2771102144	385000.0	3	3.25	1320	1327	2.0	0	0	3	8	1040	280	2008
21592	3438501320	295000.0	2	2.50	1630	1368	2.0	0	0	3	7	1280	350	2009
21593	7853361370	555000.0	4	2.50	3310	6500	2.0	0	0	3	8	3310	0	2012
21594	8564860280	459990.0	3	2.50	2680	5539	2.0	0	0	3	8	2680	0	2013
21595	123059127	625000.0	4	3.25	2730	54014	1.0	0	0	3	9	1560	1170	2007
21596	3832050580	300000.0	3	2.50	2540	5050	2.0	0	0	3	7	2540	0	2006

Duplicate rows

Most frequently occurring

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_basement	yr_built
0	1825069031	550000.0	4	1.75	2410	8447	2.0	0	3	4	8	2060	350	1936
1	6308000010	585000.0	3	2.50	2290	5089	2.0	0	0	3	9	2290	0	2001
2	8648900110	555000.0	3	2.50	1940	3211	2.0	0	0	3	8	1940	0	2009

Splitting X into :- Test & Split

```
In [16]: from sklearn.model_selection import train_test_split  
train_X, test_X, train_y, test_y = train_test_split(X, y, test_size = 0.2, random_state = 2)
```

Implementing Linear Regression

```
In [17]: from sklearn.linear_model import LinearRegression  
lin_reg = LinearRegression()  
lin_reg.fit(train_X, train_y)
```

```
Out[17]: LinearRegression()
```

Predicting Results

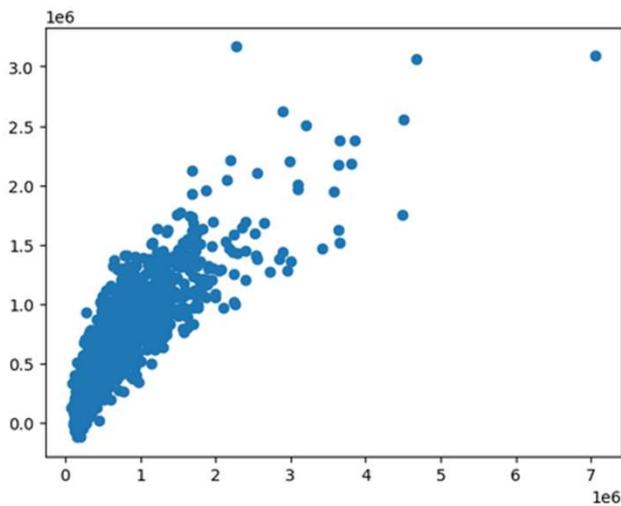
```
In [18]: y_pred = lin_reg.predict(test_X)
```

```
In [19]: y_pred
```

```
Out[19]: array([287427.00051628, 453058.91054017, 476005.88972716, ...,  
235061.9013533 , 491208.01432038, 396778.71761236])
```

Calculating MAPE(Mean Absolute Percentage Error)

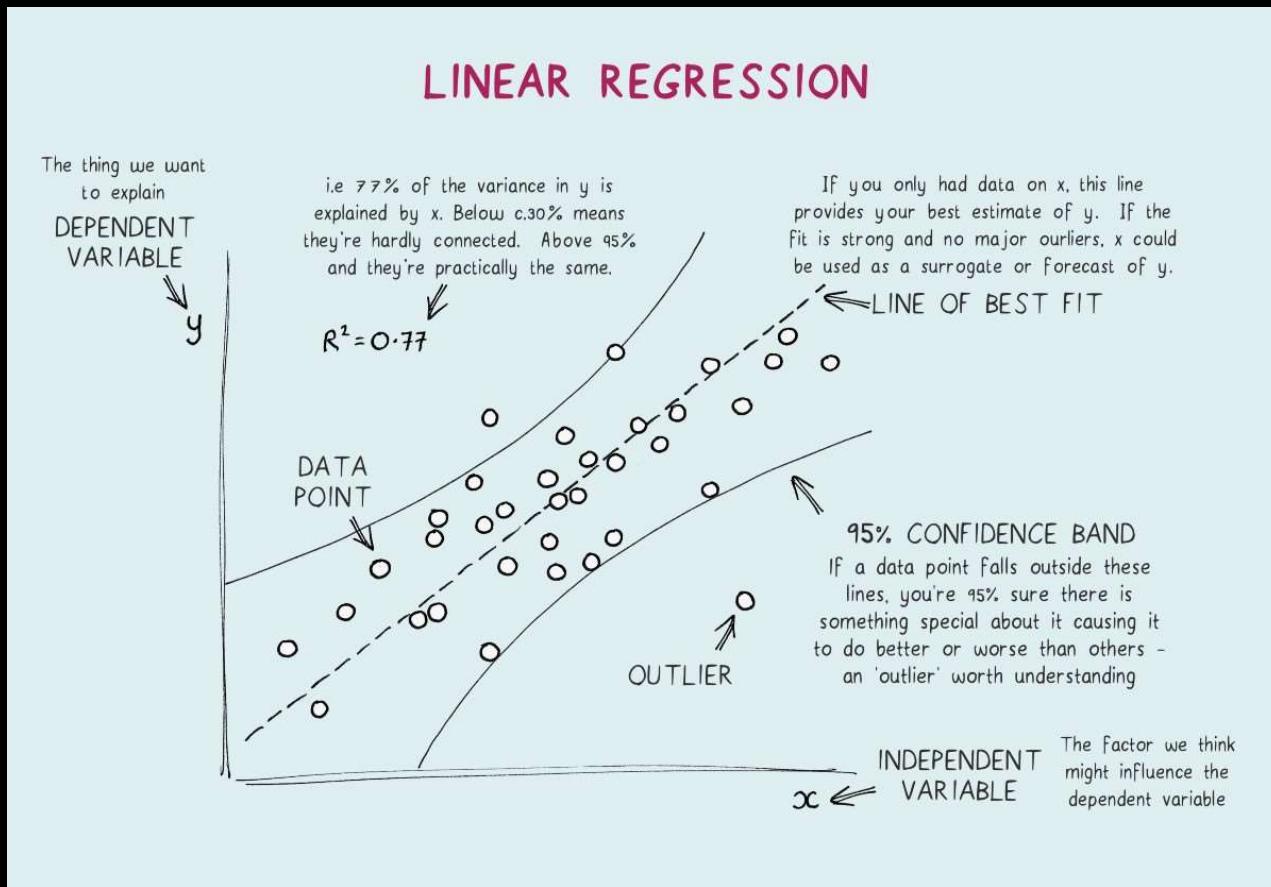
```
In [20]: from sklearn.metrics import mean_absolute_percentage_error  
mean_absolute_percentage_error(test_y, y_pred)*100  
Out[20]: 25.102180232553984  
  
In [21]: # Pretty Good Scores....That means Our MAPE is 25.10 %  
# or we can say that our predicted values are "25.10 % deviated" from our Actual Results  
  
In [22]: # we can easily calculate Accuracy of our Linear Regression Model By Subtracting (100-MAPE)%  
  
In [23]: 100 - 25.10  
Out[23]: 74.9  
  
In [24]: plt.scatter(test_y, y_pred)  
Out[24]: <matplotlib.collections.PathCollection at 0x1eeb7a9bac0>
```



We Have Got 74.9% of Accuracy

(Using Scikit Learn Library & With Outliers)

LINEAR REGRESSION (AFTER TREATING OUTLIERS)

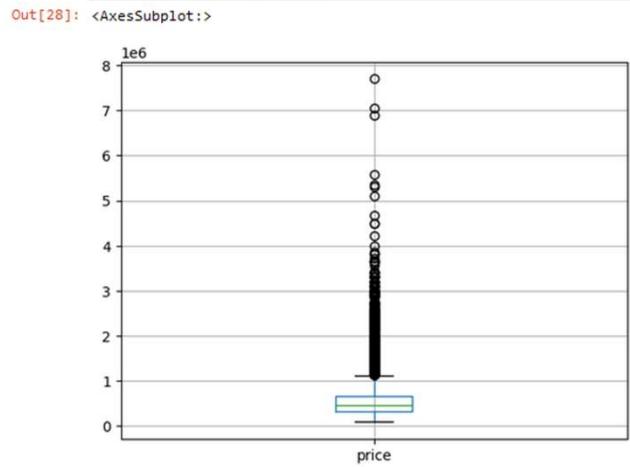


Now Doing it after removing Outliers

```
In [25]: # we will first create a copy of our dataset-"data", and create a DataFrame-"df"
In [26]: df = data.copy()
In [27]: df.head()
```

```
Out[27]:
      id      price  bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront  view  condition  grade  sqft_above
0  9106000005  1310000.0        4       2.25     3750    5000    2.0          0     0     5     8     2440
1  5101400871   445500.0        2       1.75     1390    6670    1.0          0     0     3     6      720
2  7923800250   450000.0        5       2.00     1870    7344    1.5          0     0     3     7     1870
3  8730000270   359000.0        2       2.75     1370    1140    2.0          0     0     3     8     1080
4  9178601860   1700000.0        5       3.00     3320    5354    2.0          0     0     3     9     3320
```

```
In [28]: #we know that to find Outliers, Boxplot is Best Visualization plot
df.boxplot(column = ['price'])
```



```
In [29]: #there is a significant amount of outliers...so we have to treat them
```

Function to Treat Outliers

```
In [30]: def OutRem (data, Var_name):
    Q1 = data[Var_name].quantile(0.25)  #first Quartile
    Q3 = data[Var_name].quantile(0.75)  #third Quartile
    IQR = Q3 - Q1                     #Calculating Inter-Quartile Range
    data = data.loc[~((data[Var_name] < (Q1 - 1.5 * IQR)) | (data[Var_name] > (Q3 + 1.5 * IQR))),]
    return data
```



```
In [31]: # Treating price's Outliers
df = OutRem(df,"price")
```



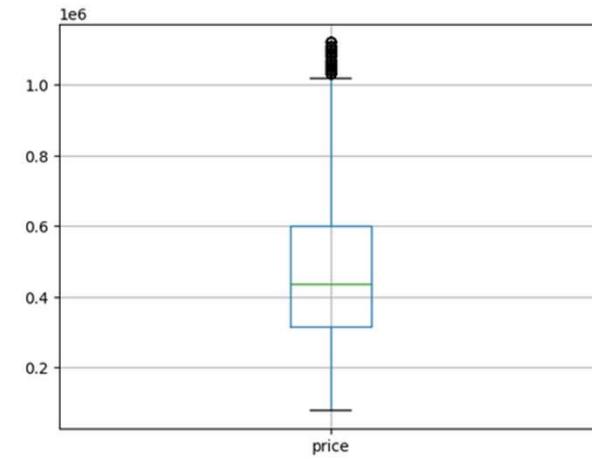
```
In [32]: %matplotlib inline
```



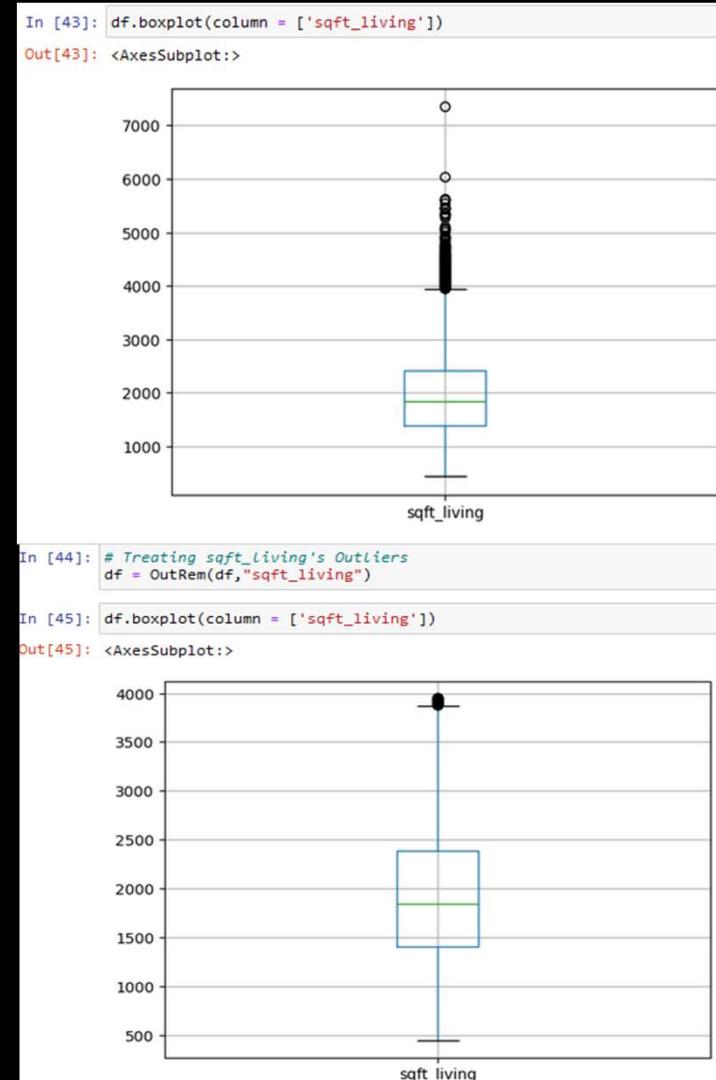
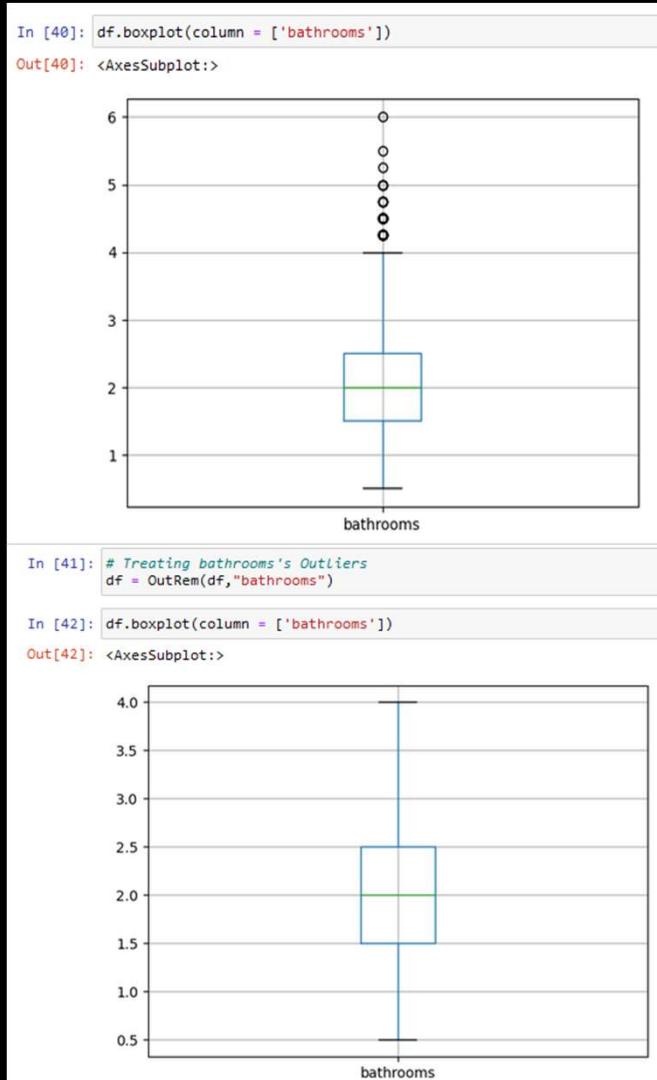
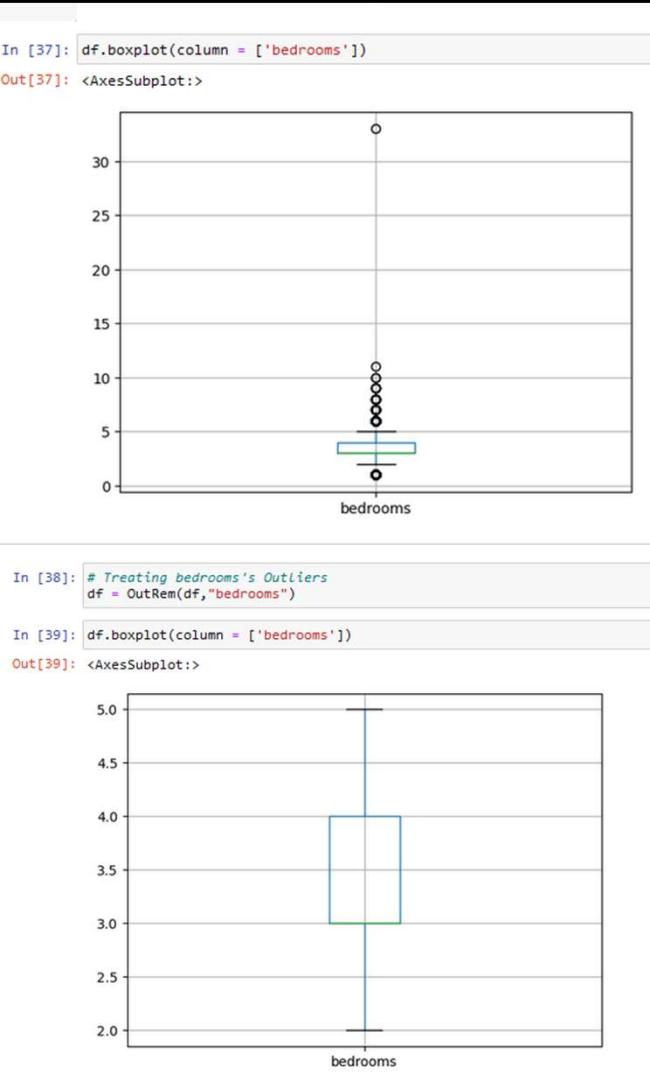
```
In [33]: #Again checking to Confirm that Outliers have been treated.
df.boxplot(column = ['price'])
```



```
Out[33]: <AxesSubplot:>
```



```
In [34]: # we can see the change clearly, outliers have been treated to an extent.
# we can ignore that small outliers still showing
```



HARSH YADAV

```
In [46]: df = OutRem(df,"sqft_lot")
In [47]: df = OutRem(df,"floors")
In [48]: df = OutRem(df,"condition")
In [49]: df = OutRem(df,"grade")
In [50]: df = OutRem(df,"sqft_above")
In [51]: df = OutRem(df,"sqft_basement")
In [52]: df = OutRem(df,"yr_built")
In [53]: df = OutRem(df,"zipcode")
df = OutRem(df,"lat")
df = OutRem(df,"long")
df = OutRem(df,"sqft_living")
df = OutRem(df,"sqft_lot15")
# I have not Included certain Variables here, because they were type of binary, Like waterfront,view etc.

In [54]: # again doing same steps, as we did with dataset-"data"
y = df.loc[:, "price"]
X = df.drop("price", axis = 1)
In [55]: #SPLitting train & test data
from sklearn.model_selection import train_test_split
train_X, test_X, train_y, test_y = train_test_split(X, y, test_size = 0.2, random_state = 2)
In [56]: # Again Linear Regression
from sklearn.linear_model import LinearRegression
lin_reg = LinearRegression()
lin_reg.fit(train_X, train_y)
Out[56]: LinearRegression()
In [57]: y_pred_wo = lin_reg.predict(test_X)
In [58]: len(y_pred_wo)
Out[58]: 3092
In [59]: len(test_y)
Out[59]: 3092
```

MAPE(Without Outliers)

```
In [60]: mean_absolute_percentage_error(test_y, y_pred_wo)*100
```

```
Out[60]: 19.781135275315087
```

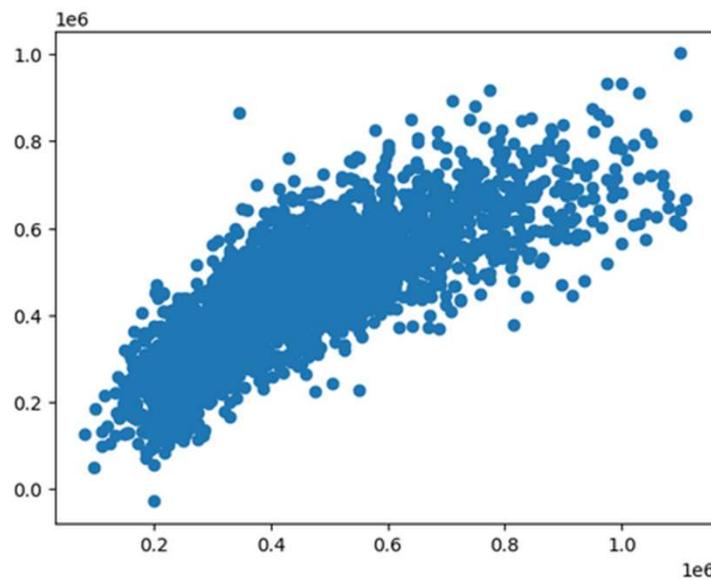
```
In [61]: # we have improvised the model by removing the outliers and by making it more balanced
```

```
In [62]: 100 - 19.78
```

```
Out[62]: 80.22
```

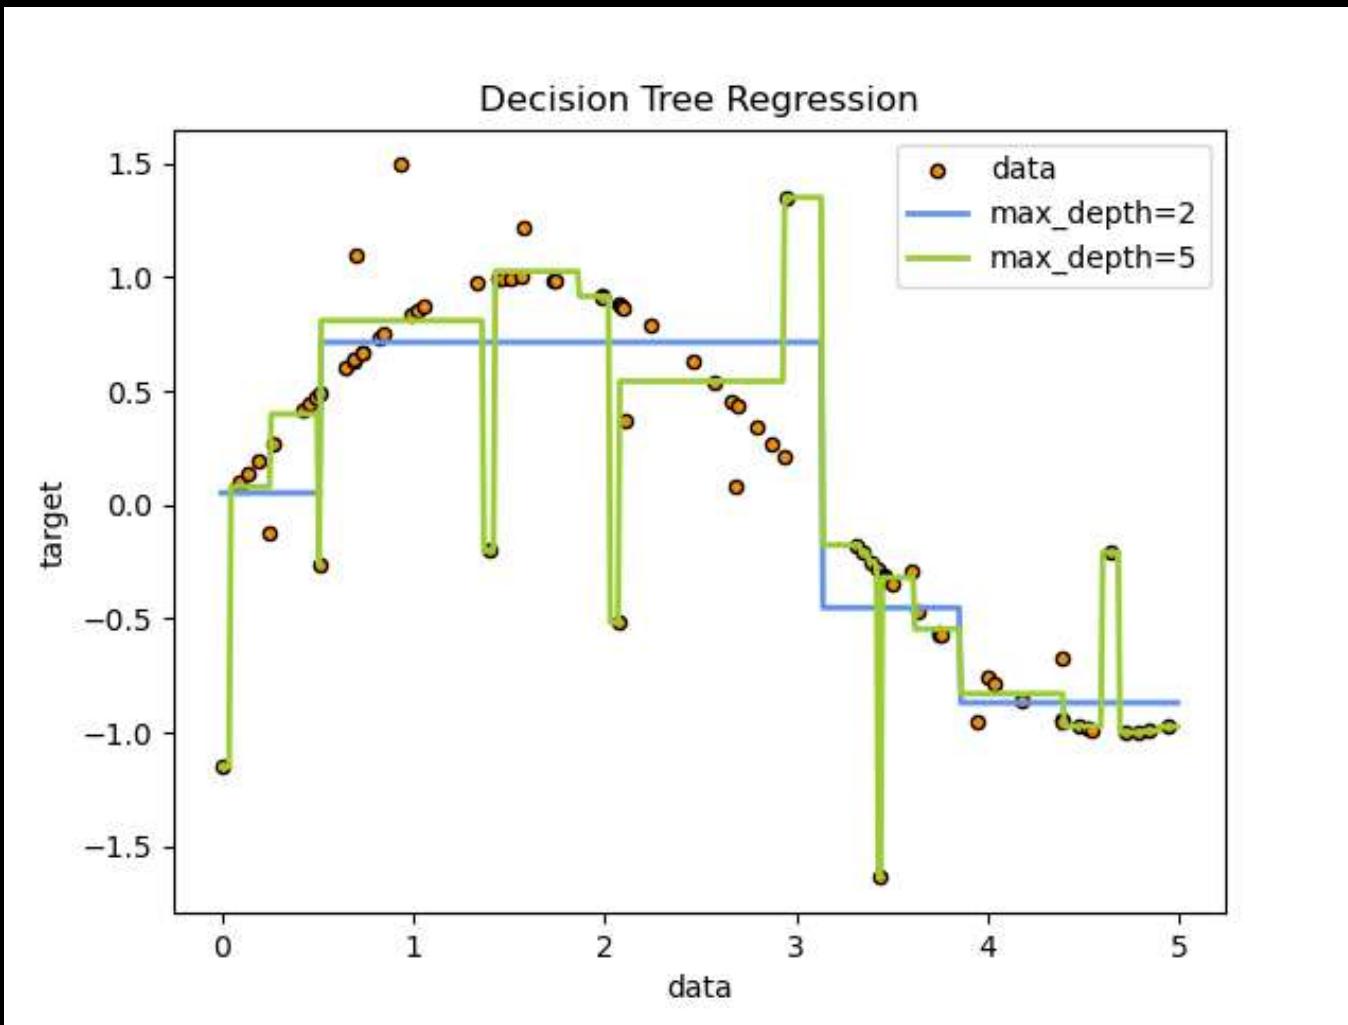
```
In [63]: plt.scatter(test_y, y_pred_wo)
```

```
Out[63]: <matplotlib.collections.PathCollection at 0xleeb7dbca60>
```



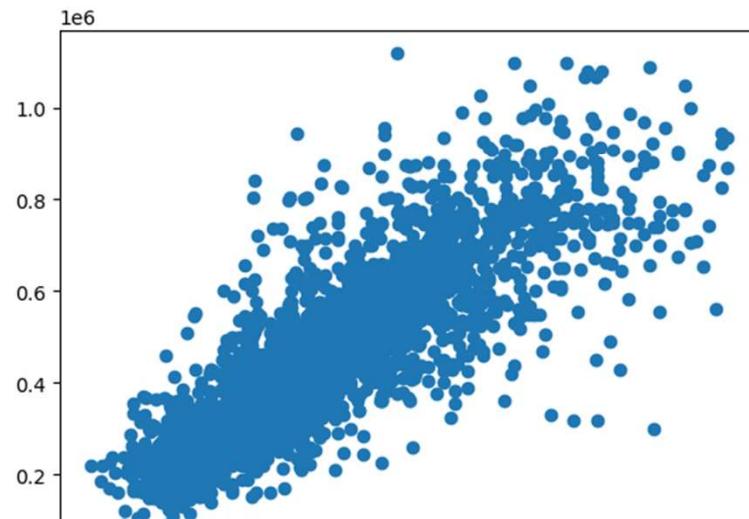
WE HAVE FINALLY ACHIEVED AN ACCURACY OF "80.22%" BY REMOVING OUTLIERS FROM SOME VARIABLES

DECISION TREE REGRESSION

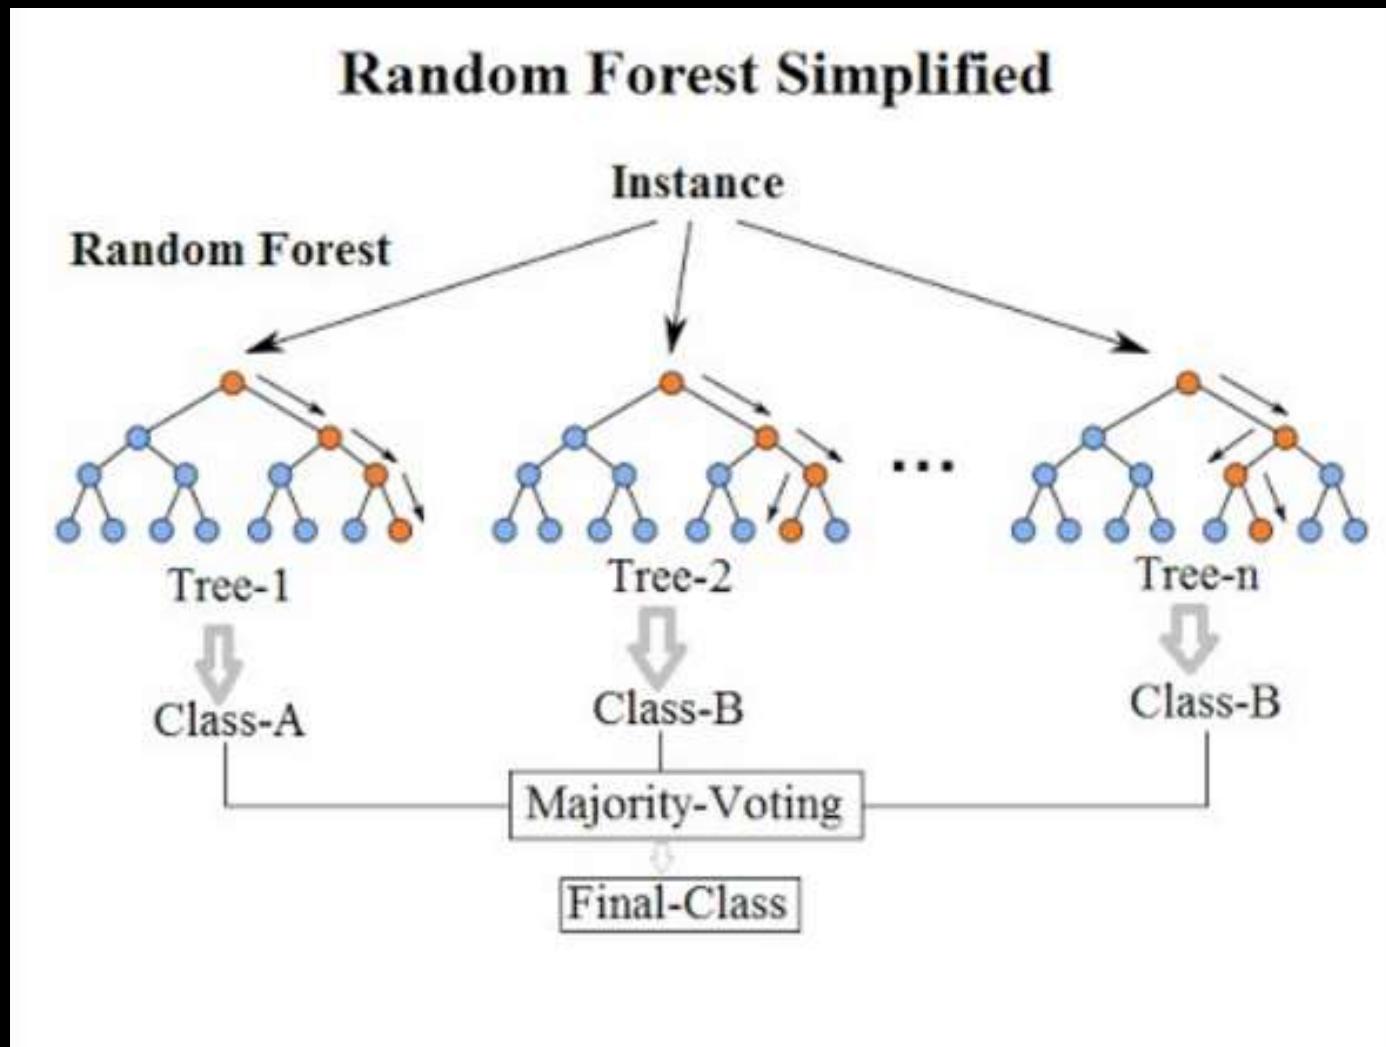


Decision Tree Regression

```
In [64]: from sklearn.tree import DecisionTreeRegressor  
  
In [65]: dt = DecisionTreeRegressor(random_state = 2)  
dt.fit(train_X, train_y)  
  
Out[65]: DecisionTreeRegressor(random_state=2)  
  
In [66]: y_pred_dt = dt.predict(test_X)  
  
In [67]: mean_absolute_percentage_error(test_y, y_pred_dt)  
  
Out[67]: 0.16507669880726836  
  
In [68]: 100-16.50  
  
Out[68]: 83.5  
  
In [69]: plt.scatter(test_y, y_pred_dt)  
  
Out[69]: <matplotlib.collections.PathCollection at 0x1eeb7f0c2e0>
```

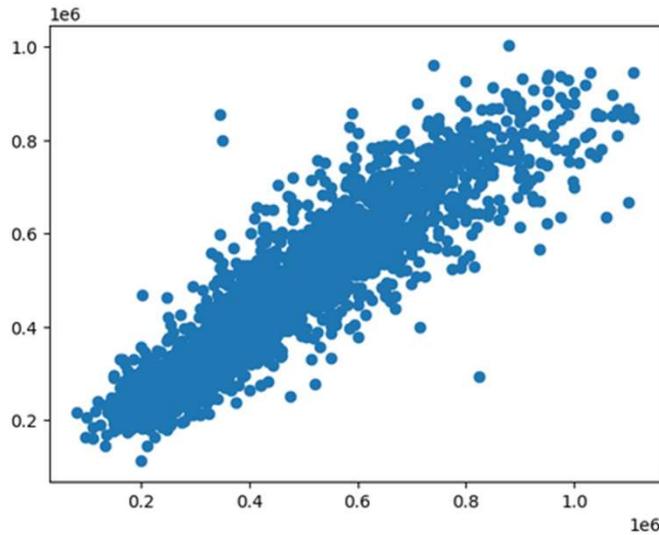


RANDOM FOREST REGRESSION

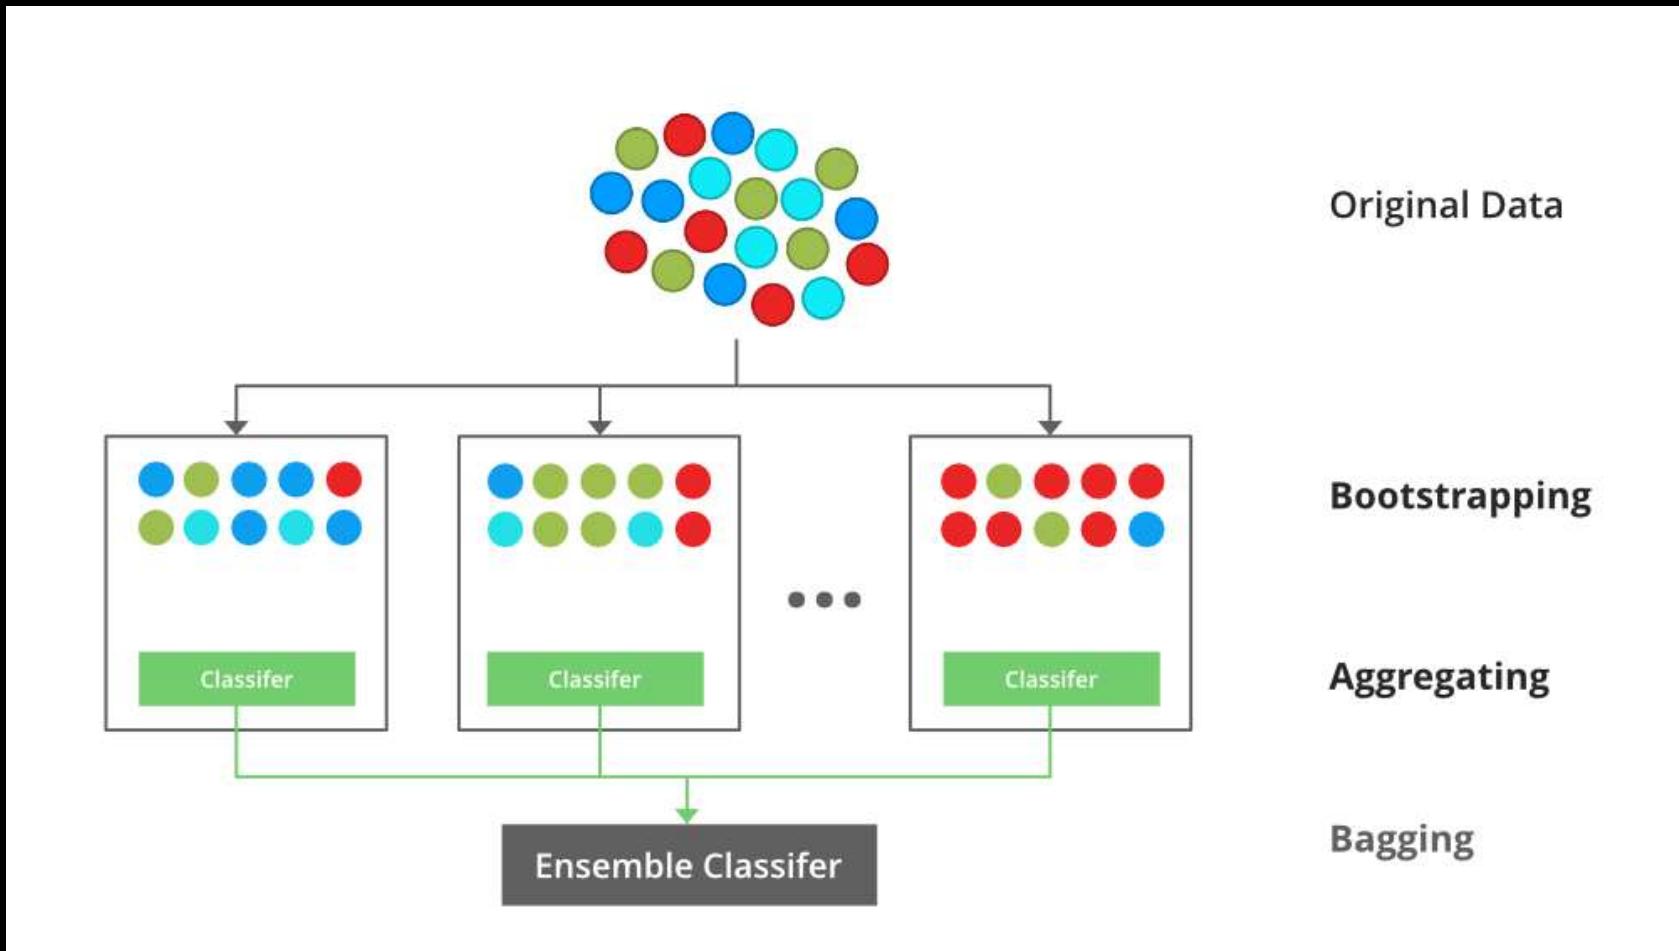


RANDOM FOREST REGRESSOR

```
In [70]: from sklearn.ensemble import RandomForestRegressor  
  
In [71]: rf = RandomForestRegressor(n_estimators = 1000, random_state = 2)  
rf.fit(train_X, train_y)  
  
Out[71]: RandomForestRegressor(n_estimators=1000, random_state=2)  
  
In [72]: y_pred_rf = rf.predict(test_X)  
  
In [73]: mean_absolute_percentage_error(test_y, y_pred_rf)  
  
Out[73]: 0.11566496080803386  
  
In [74]: 100-11.56  
  
Out[74]: 88.44  
  
In [75]: plt.scatter(test_y, y_pred_rf)  
Out[75]: <matplotlib.collections.PathCollection at 0x1eabb8ca7f0>
```



XGBOOST



XGboost

```
In [86]: !pip install xgboost
```

```
Requirement already satisfied: xgboost in c:\users\anika\anaconda3\lib\site-packages (1.7.1)
Requirement already satisfied: scipy in c:\users\anika\anaconda3\lib\site-packages (from xgboost) (1.9.1)
Requirement already satisfied: numpy in c:\users\anika\anaconda3\lib\site-packages (from xgboost) (1.21.5)
```

```
In [87]: import xgboost as xg
```

```
In [88]: xgb_r = xg.XGBRFRegressor()
xgb_r.fit(train_X, train_y)
```

```
Out[88]: XGBRFRegressor(base_score=0.5, booster='gbtree', callbacks=None,
                       colsample_bylevel=1, colsample_bytree=1,
                       early_stopping_rounds=None, enable_categorical=False,
                       eval_metric=None, feature_types=None, gamma=0, gpu_id=-1,
                       grow_policy='depthwise', importance_type=None,
                       interaction_constraints='', max_bin=256, max_cat_threshold=64,
                       max_cat_to_onehot=4, max_delta_step=0, max_depth=6, max_leaves=0,
                       min_child_weight=1, missing=nan, monotone_constraints='()', n_estimators=100, n_jobs=0, num_parallel_tree=100,
                       objective='reg:squarederror', predictor='auto', random_state=0,
                       reg_alpha=0, ...)
```

```
In [89]: y_pred_x = xgb_r.predict(test_X)
```

```
In [90]: mean_absolute_percentage_error(test_y, y_pred_x)
```

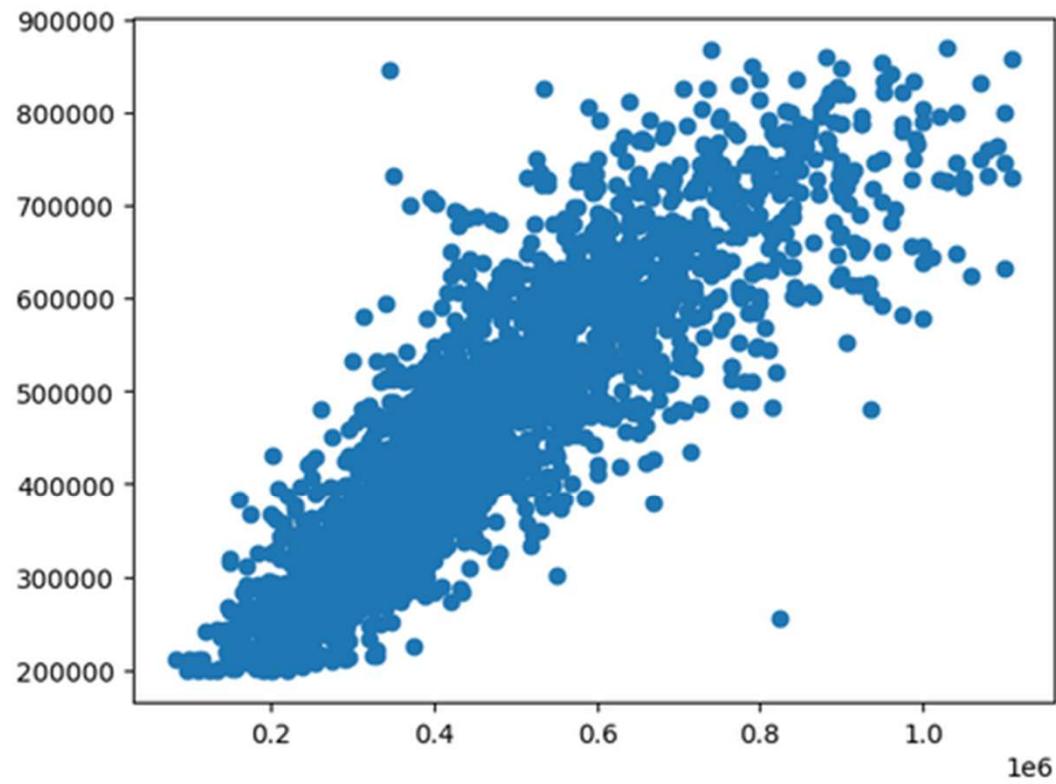
```
Out[90]: 0.14498487241580113
```

```
In [91]: 100-14.49
```

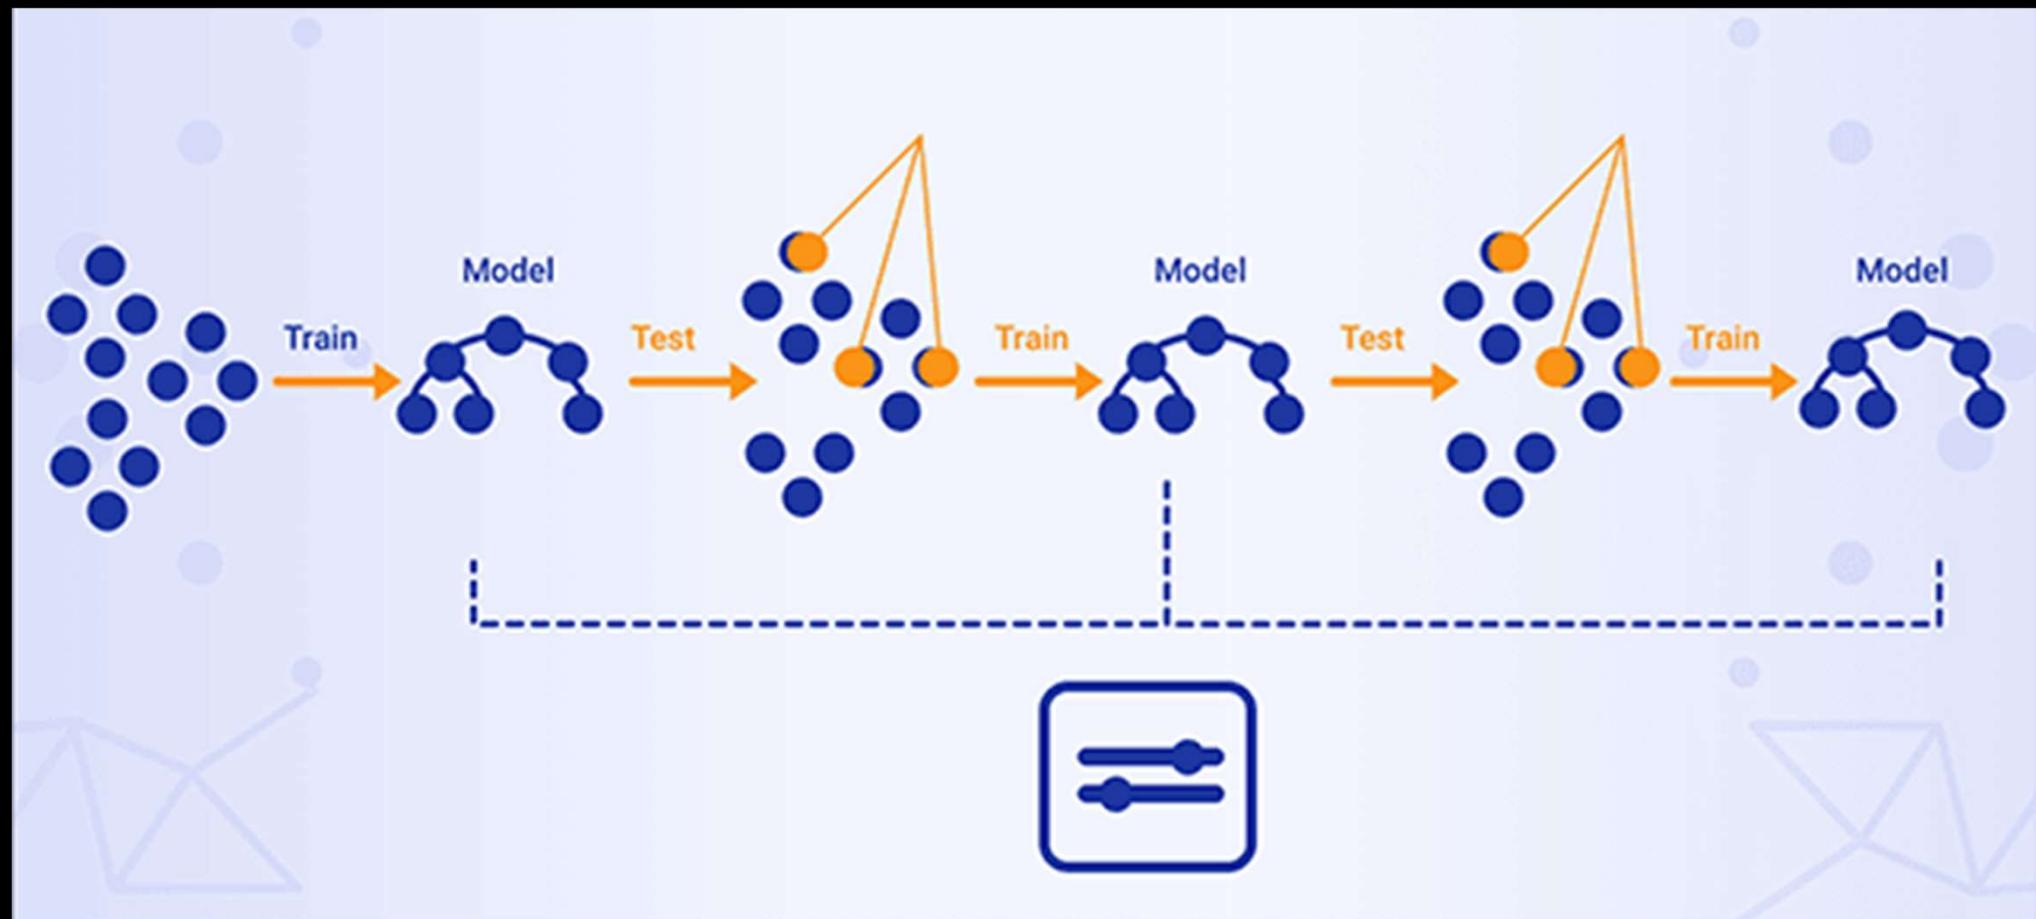
```
Out[91]: 85.51
```

```
In [92]: plt.scatter(test_y, y_pred_x)
```

```
Out[92]: <matplotlib.collections.PathCollection at 0x270b18bdc70>
```



GRADIENT BOOSTING REGRESSION



HARSH YADAV

GRADIENT BOOSTING REGRESSION

```
In [77]: from sklearn.ensemble import GradientBoostingRegressor  
clf = GradientBoostingRegressor(n_estimators = 400, max_depth = 5, min_samples_split = 2, learning_rate = 0.1, loss =
```

n_estimator — The number of boosting stages to perform. We should not set it too high which would overfit our model

max_depth — The depth of the tree node

learning_rate — Rate of learning the data

loss — loss function

minimum sample split — Number of sample to be split for learning the data

```
In [78]: clf.fit(train_X, train_y)
```

```
Out[78]: GradientBoostingRegressor(max_depth=5, n_estimators=400)
```

```
In [79]: y_pred_clf = clf.predict(test_X)
```

```
In [80]: mean_absolute_percentage_error(test_y, y_pred_clf)*100
```

```
Out[80]: 10.99443490267817
```

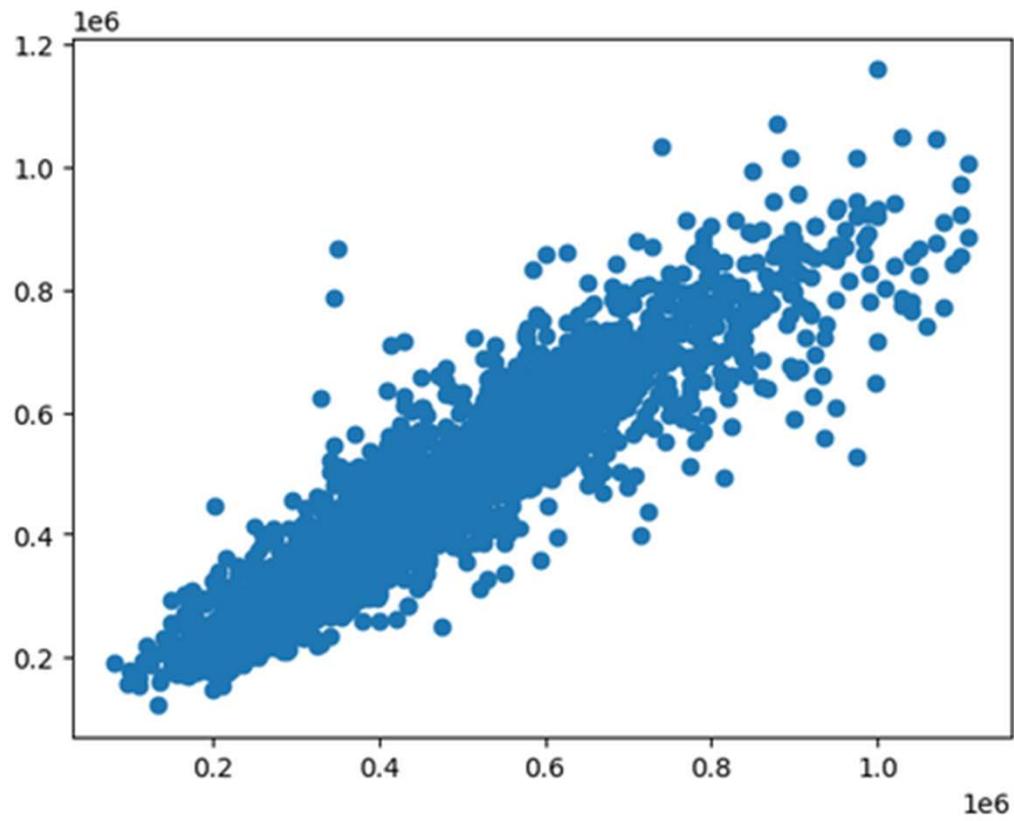
```
In [81]: #MAPE is 10.99  
100-10.99
```

```
Out[81]: 89.01
```

```
In [82]: # point to be noted : the accuracy is between - ( test"y" & predicted"y" )
```

```
In [83]: plt.scatter(test_y, y_pred_clf)
```

```
Out[83]: <matplotlib.collections.PathCollection at 0x1eefb3d9460>
```



We Have Achieved the Accuracy of "89.01%" by GRADIENT BOOSTING REGRESSION.

```
In [100]: # model.score automates the prediction of the data using X_test and compares it with Y_test and by  
# default uses the R-squared metric to do so (hence don't need to manually derive y_pred).  
clf.score(test_X, test_y)
```

```
Out[100]: 0.8659642506327195
```

```
In [103]: #we can also check this using r2_score  
sklearn.metrics.r2_score(test_y, y_pred_clf)
```

```
Out[103]: 0.8659642506327195
```

CROSS-VALIDATION SCORES OF SOME MODELS

```
In [109]: from sklearn.model_selection import ShuffleSplit  
from sklearn.model_selection import cross_val_score
```

```
In [110]: # cv_score of "RandomForestRegressor"  
cv = ShuffleSplit(n_splits=5, test_size=0.2, random_state=0)  
cross_val_score(RandomForestRegressor(), X, y, cv=cv)
```

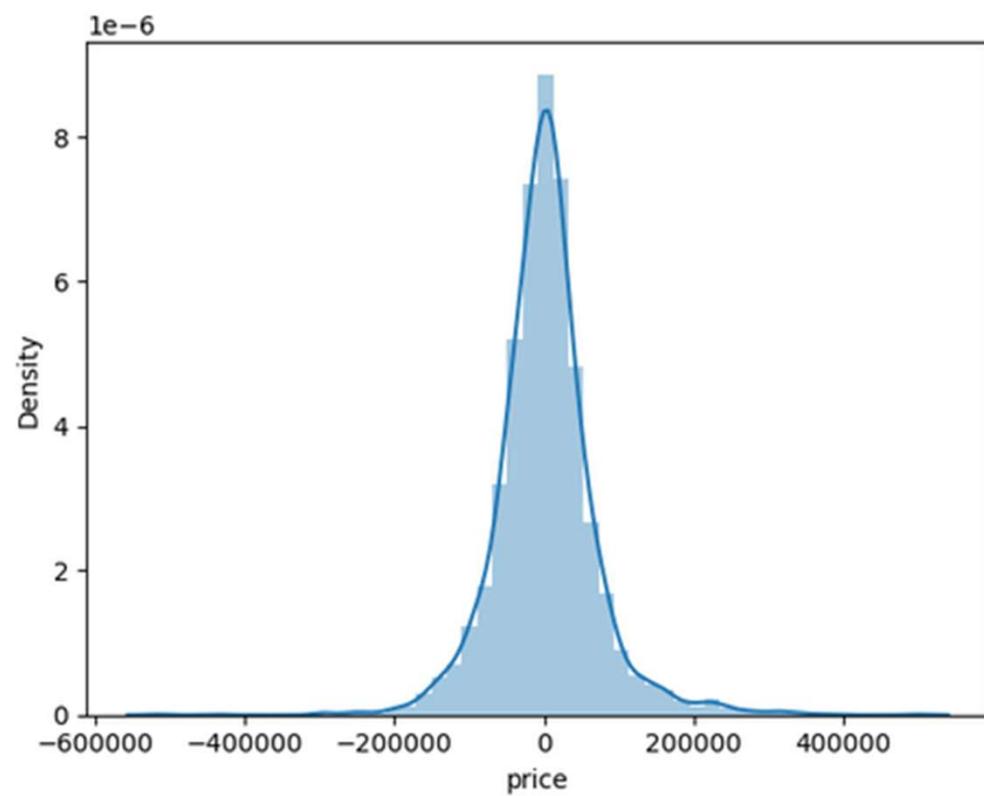
```
Out[110]: array([0.85870483, 0.85035836, 0.85279999, 0.83998357, 0.8604371 ])
```

```
In [ ]:
```

```
In [111]: # cv_score of "GradientBoostingRegressor"  
cv = ShuffleSplit(n_splits=5, test_size=0.2, random_state=0)  
cross_val_score(GradientBoostingRegressor(), X, y, cv=cv)
```

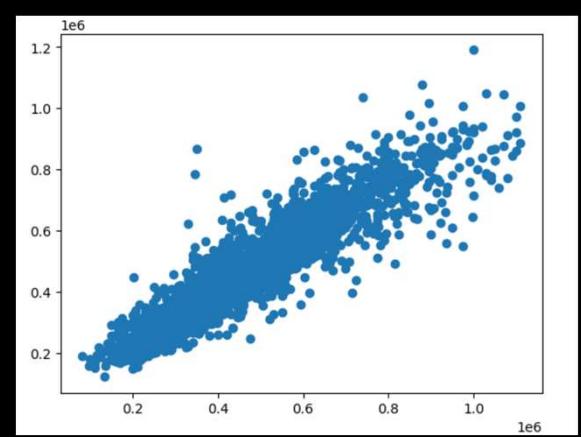
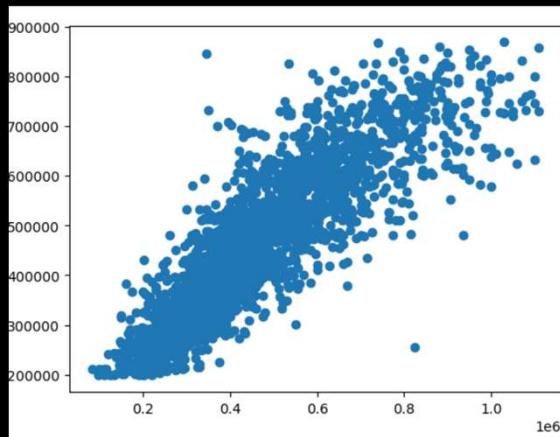
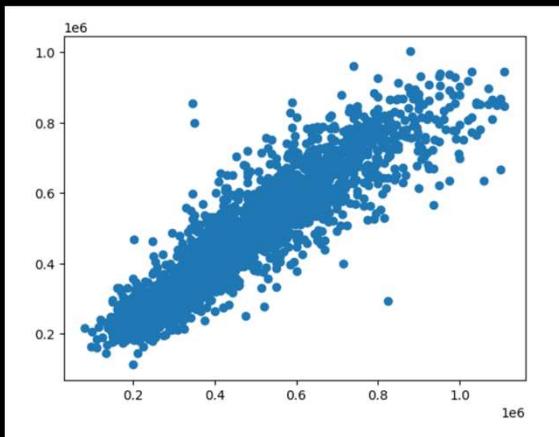
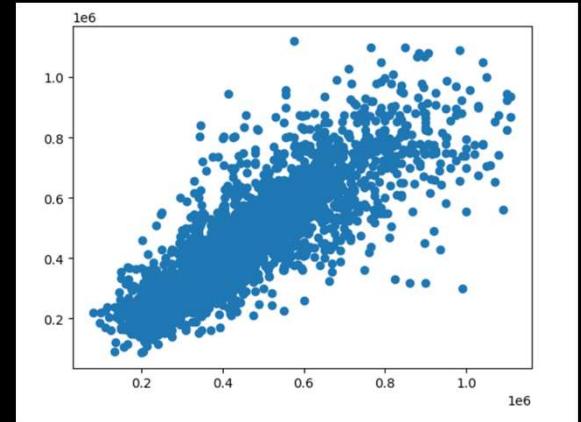
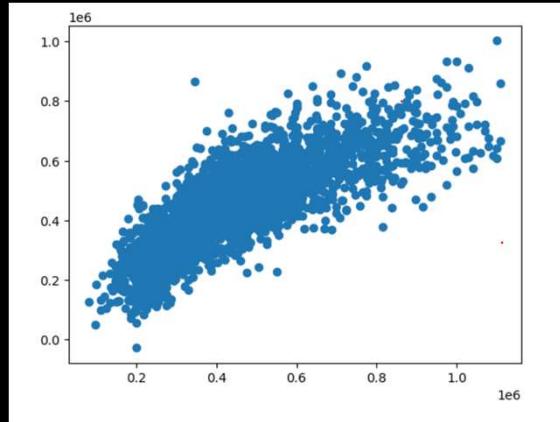
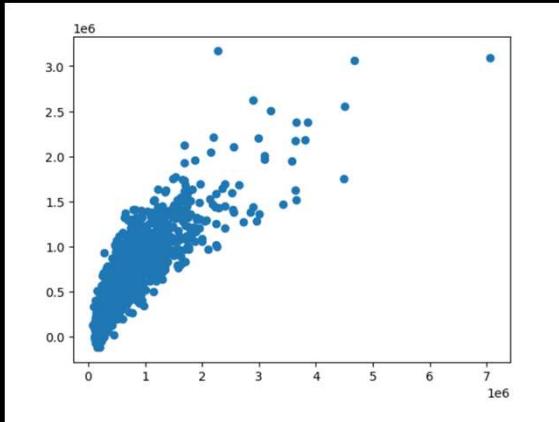
```
Out[111]: array([0.84148168, 0.83184076, 0.84058511, 0.82753675, 0.84442999])
```

```
In [120]: sns.distplot((test_y-y_pred_clf),bins=50);# to check that there is no skewness & our predictions have been computed bal  
C:\Users\anika\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)
```



HARSH YADAV

COMPARISON BETWEEN PREDICTED RECORDS AND ORIGINAL RECORDS BY DIFFERENT TECHNIQUES WE'VE USED



THE END
THANK YOU

SUBMITTED BY – HARSH YADAV

SUBMITTED TO – DHRUV
GOVIL SIR

HARSH YADAV